

supervised-car-pricing-model

November 27, 2023

1 Project Name - Car Pricing Model

Project Type - EDA/Regression

Contribution - Individual

2 Project Summary -

The Chinese automobile company Geely Auto aimed to enter the US market and hired consultants to analyze the American car pricing dynamics through a data-driven regression model in order to optimally design and price their vehicles. Multiple regression techniques were applied, among which the linear regression model with the highest accuracy and generalizability established a quantifiable relationship between car features like horsepower, weight, engine-size and pricing. It enabled Geely to determine target specification ranges based on competitive analysis for given desired price points. The model provides data-backed guidance to strategically optimize their automobile configurations and dynamic pricing strategy in order to successfully penetrate the unfamiliar American car market. With continuous model re-training using latest market data, Geely can account for evolving pricing trends and adapt their pricing accordingly over time. Overall, the fitted linear model supplies actionable insights for data-informed decision making regarding optimal vehicle design and pricing to assist Geely's entry and viability in the US automobile industry.

3 GitHub Link -

<https://github.com/shantanu0101/Car-Pricing-Model-Regression-Model->

4 Problem Statement

A Chinese automobile company Geely Auto aspires to enter the US market by setting up their manufacturing unit there and producing cars locally to give competition to their US and European counterparts.

They have contracted an automobile consulting company to understand the factors on which the pricing of cars depends. Specifically, they want to understand the factors affecting the pricing of cars in the American market, since those may be very different from the Chinese market. The company wants to know:

Which variables are significant in predicting the price of a car
How well those variables describe the price of a car
Based on various market surveys, the consulting firm has gathered a large data

set of different types of cars across the America market.

5 General Guidelines : -

1. Well-structured, formatted, and commented code is required.
2. Exception Handling, Production Grade Code & Deployment Ready Code will be a plus. Those students will be awarded some additional credits.

The additional credits will have advantages over other students during Star Student selection.

[Note: - Deployment Ready Code is defined as, the whole .ipynb notebook should be executable without a single error logged.]

3. Each and every logic should have proper comments.
4. You may add as many number of charts you want. Make Sure for each and every chart the following format should be answered.

Chart visualization code

- Why did you pick the specific chart?
- What is/are the insight(s) found from the chart?
- Will the gained insights help creating a positive business impact? Are there any insights that lead to negative growth? Justify with specific reason.

[Hints : - Do the Vizualization in a structured way while following “UBM” Rule.

U - Univariate Analysis,

B - Bivariate Analysis (Numerical - Categorical, Numerical - Numerical, Categorical - Categorical)

M - Multivariate Analysis]

6. You may add more ml algorithms for model creation. Make sure for each and every algorithm, the following format should be answered.
 - Explain the ML Model used and it's performance using Evaluation metric Score Chart.
 - Cross- Validation & Hyperparameter Tuning
 - Explain each evaluation metric's indication towards business and the business impact pf the ML model used.

6 *Let's Begin !*

6.1 *1. Know Your Data*

6.1.1 Import Libraries

```
[111]: # Import Libraries
import math # Import the math module for mathematical operations
import numpy as np # Import NumPy for numerical operations
import pandas as pd # Import Pandas for data manipulation
import seaborn as sns # Import Seaborn for statistical data visualization
```

```

import matplotlib.pyplot as plt # Import Matplotlib for plotting

from sklearn.preprocessing import StandardScaler # Import StandardScaler for
↳standardization of features
from sklearn.preprocessing import MinMaxScaler # Import MinMaxScaler for
↳scaling features to a range
from sklearn.metrics import mean_squared_error # Import mean_squared_error for
↳calculating Mean Squared Error
from sklearn.metrics import mean_absolute_error # Import mean_absolute_error
↳for calculating Mean Absolute Error
from sklearn.metrics import mean_absolute_percentage_error # Import
↳mean_absolute_percentage_error for calculating MAPE
from sklearn.metrics import r2_score # Import r2_score for calculating
↳R-squared

from scipy.stats import pointbiserialr # Import pointbiserialr for
↳point-biserial correlation coefficient
from sklearn.model_selection import train_test_split # Import train_test_split
↳for splitting data into train and test sets
from sklearn.model_selection import cross_val_score, cross_val_predict #
↳Import cross-validation functions
from sklearn.model_selection import GridSearchCV # Import GridSearchCV for
↳hyperparameter tuning
from sklearn.linear_model import LinearRegression # Import LinearRegression
↳model
from sklearn.linear_model import Ridge # Import Ridge Regression model
from sklearn.linear_model import Lasso # Import Lasso Regression model
from sklearn.linear_model import ElasticNet # Import ElasticNet model

```

6.1.2 Dataset Loading

```

[112]: # Load Dataset
from google.colab import drive
drive.mount('/content/drive')

```

Drive already mounted at /content/drive; to attempt to forcibly remount, call `drive.mount("/content/drive", force_remount=True)`.

```

[113]: path = '/content/drive/MyDrive/CarPrice_project.csv'
cars_df = pd.read_csv(path)

```

6.1.3 Dataset First View

```
[114]: # Dataset First Look
cars_df
```

```
[114]:
```

	car_ID	symboling	CarName	fueltype	aspiration	\
0	1	3	alfa-romero giulia	gas	std	
1	2	3	alfa-romero stelvio	gas	std	
2	3	1	alfa-romero Quadrifoglio	gas	std	
3	4	2	audi 100 ls	gas	std	
4	5	2	audi 100ls	gas	std	
..	
200	201	-1	volvo 145e (sw)	gas	std	
201	202	-1	volvo 144ea	gas	turbo	
202	203	-1	volvo 244dl	gas	std	
203	204	-1	volvo 246	diesel	turbo	
204	205	-1	volvo 264gl	gas	turbo	

	doornumber	carbody	drivewheel	engine	location	wheelbase	...	\
0	two	convertible	rwd	front	88.6	...		
1	two	convertible	rwd	front	88.6	...		
2	two	hatchback	rwd	front	94.5	...		
3	four	sedan	fwd	front	99.8	...		
4	four	sedan	4wd	front	99.4	...		
..		
200	four	sedan	rwd	front	109.1	...		
201	four	sedan	rwd	front	109.1	...		
202	four	sedan	rwd	front	109.1	...		
203	four	sedan	rwd	front	109.1	...		
204	four	sedan	rwd	front	109.1	...		

	enginesize	fuelsystem	boreratio	stroke	compressionratio	horsepower	\
0	130	mpfi	3.47	2.68	9.0	111	
1	130	mpfi	3.47	2.68	9.0	111	
2	152	mpfi	2.68	3.47	9.0	154	
3	109	mpfi	3.19	3.40	10.0	102	
4	136	mpfi	3.19	3.40	8.0	115	
..	
200	141	mpfi	3.78	3.15	9.5	114	
201	141	mpfi	3.78	3.15	8.7	160	
202	173	mpfi	3.58	2.87	8.8	134	
203	145	idi	3.01	3.40	23.0	106	
204	141	mpfi	3.78	3.15	9.5	114	

	peakrpm	citympg	highwaympg	price
0	5000	21	27	13495.0
1	5000	21	27	16500.0

2	5000	19	26	16500.0
3	5500	24	30	13950.0
4	5500	18	22	17450.0
..
200	5400	23	28	16845.0
201	5300	19	25	19045.0
202	5500	18	23	21485.0
203	4800	26	27	22470.0
204	5400	19	25	22625.0

[205 rows x 26 columns]

6.1.4 Dataset Rows & Columns count

```
[115]: # Dataset Rows & Columns count
cars_df.shape
```

```
[115]: (205, 26)
```

6.1.5 Dataset Information

```
[116]: # Dataset Info
cars_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 205 entries, 0 to 204
Data columns (total 26 columns):
#   Column                Non-Null Count  Dtype
---  -
0   car_ID                 205 non-null    int64
1   symboling              205 non-null    int64
2   CarName                205 non-null    object
3   fueltype               205 non-null    object
4   aspiration              205 non-null    object
5   doornumber             205 non-null    object
6   carbody                205 non-null    object
7   drivewheel             205 non-null    object
8   enginelocation         205 non-null    object
9   wheelbase              205 non-null    float64
10  carlength              205 non-null    float64
11  carwidth               205 non-null    float64
12  carheight              205 non-null    float64
13  curbweight             205 non-null    int64
14  enginetype             205 non-null    object
15  cylindernumber         205 non-null    object
16  enginesize             205 non-null    int64
17  fuelsystem             205 non-null    object
```

```

18  boreratio          205 non-null    float64
19  stroke             205 non-null    float64
20  compressionratio   205 non-null    float64
21  horsepower         205 non-null    int64
22  peakrpm            205 non-null    int64
23  citympg            205 non-null    int64
24  highwaympg         205 non-null    int64
25  price              205 non-null    float64
dtypes: float64(8), int64(8), object(10)
memory usage: 41.8+ KB

```

Duplicate Values

```
[117]: # Dataset Duplicate Value Count
cars_df.duplicated().sum()
```

```
[117]: 0
```

Missing Values/Null Values


```
[118]: # Missing Values/Null Values Count
cars_df.isnull().sum()
```

```
[118]: car_ID          0
symboling          0
CarName            0
fueltype           0
aspiration         0
doornumber         0
carbody            0
drivewheel         0
enginelocation     0
wheelbase          0
carlength          0
carwidth           0
carheight          0
curbweight         0
enginetype         0
cylindernumber     0
enginesize         0
fuelsystem         0
boreratio          0
stroke             0
compressionratio   0
horsepower         0
peakrpm            0
citympg            0
highwaympg         0

```

```
price
dtype: int64
```

```
[119]: # Visualizing the missing values
plt.figure(figsize = (8,6), dpi = 100)
sns.heatmap(cars_df.isnull(), cmap = 'viridis', cbar = False, yticklabels =
↪False)
plt.show()
```



6.1.6 What did you know about your dataset?

The dataset comprises information on various attributes related to automobiles, encompassing both categorical and numerical features. Each entry in the dataset is associated with a unique car identification (car_ID). The features include assessments of risk (symboling), car name (CarName),

fuel type (fueltype), aspiration type (aspiration), the number of doors (doornumber), car body style (carbody), drivetrain type (drivewheel), engine location (enginelocation), and dimensions such as wheelbase, car length, width, and height. Other essential characteristics encompass curb weight, engine type, cylinder count, engine size, fuel injection system type, bore ratio, stroke, compression ratio, horsepower, peak revolutions per minute, and fuel efficiency measured in miles per gallon for both city and highway driving. The dataset culminates in the target variable, 'Price,' representing the car's cost. With these diverse features, the dataset is well-structured for predictive modeling, aiming to establish relationships between the car attributes and their corresponding prices.

6.2 2. Understanding Your Variables

```
[120]: # Dataset Columns
column_list = list(cars_df.columns)
column_list
```

```
[120]: ['car_ID',
       'symboling',
       'CarName',
       'fueltype',
       'aspiration',
       'doornumber',
       'carbody',
       'drivewheel',
       'enginelocation',
       'wheelbase',
       'carlength',
       'carwidth',
       'carheight',
       'curbweight',
       'enginetype',
       'cylindernumber',
       'enginesize',
       'fuelsystem',
       'boreratio',
       'stroke',
       'compressionratio',
       'horsepower',
       'peakrpm',
       'citympg',
       'highwaympg',
       'price']
```

```
[121]: # Dataset Describe
cars_df.describe()
```

```
[121]:
```

	car_ID	symboling	wheelbase	carlength	carwidth	carheight	\
count	205.000000	205.000000	205.000000	205.000000	205.000000	205.000000	

mean	103.000000	0.834146	98.756585	174.049268	65.907805	53.724878
std	59.322565	1.245307	6.021776	12.337289	2.145204	2.443522
min	1.000000	-2.000000	86.600000	141.100000	60.300000	47.800000
25%	52.000000	0.000000	94.500000	166.300000	64.100000	52.000000
50%	103.000000	1.000000	97.000000	173.200000	65.500000	54.100000
75%	154.000000	2.000000	102.400000	183.100000	66.900000	55.500000
max	205.000000	3.000000	120.900000	208.100000	72.300000	59.800000

	curbweight	enginesize	boreratio	stroke	compressionratio	\
count	205.000000	205.000000	205.000000	205.000000	205.000000	
mean	2555.565854	126.907317	3.329756	3.255415	10.142537	
std	520.680204	41.642693	0.270844	0.313597	3.972040	
min	1488.000000	61.000000	2.540000	2.070000	7.000000	
25%	2145.000000	97.000000	3.150000	3.110000	8.600000	
50%	2414.000000	120.000000	3.310000	3.290000	9.000000	
75%	2935.000000	141.000000	3.580000	3.410000	9.400000	
max	4066.000000	326.000000	3.940000	4.170000	23.000000	

	horsepower	peakrpm	citympg	highwaympg	price
count	205.000000	205.000000	205.000000	205.000000	205.000000
mean	104.117073	5125.121951	25.219512	30.751220	13276.710571
std	39.544167	476.985643	6.542142	6.886443	7988.852332
min	48.000000	4150.000000	13.000000	16.000000	5118.000000
25%	70.000000	4800.000000	19.000000	25.000000	7788.000000
50%	95.000000	5200.000000	24.000000	30.000000	10295.000000
75%	116.000000	5500.000000	30.000000	34.000000	16503.000000
max	288.000000	6600.000000	49.000000	54.000000	45400.000000

6.2.1 Variables Description

1. **car_ID:** Unique identifier for each car.
2. **symboling:** Risk rating associated with the car.
3. **CarName:** Name of the car.
4. **fueltype:** Type of fuel the car uses (e.g., gas or diesel).
5. **aspiration:** Type of aspiration (e.g., std or turbo).
6. **doornumber:** Number of doors on the car.
7. **carbody:** Body style of the car.
8. **drivewheel:** Type of drivetrain (e.g., front-wheel-drive, rear-wheel-drive, 9.or four-wheel-drive).
9. **enginelocation:** Location of the car engine (front or rear).
10. **wheelbase:** Distance between the centers of the front and rear wheels.
11. **carlength:** Length of the car.

12. **carwidth:** Width of the car.
13. **carheight:** Height of the car.
14. **curbweight:** Weight of the car without occupants or baggage.
15. **enginetype:** Type of engine.
16. **cylindernumber:** Number of cylinders in the engine.
17. **enginesize:** Size of the car's engine.
18. **fuelsystem:** Type of fuel injection system.
19. **boreratio:** Bore ratio of the engine.
20. **stroke:** Stroke or volume inside the engine.
21. **compressionratio:** Compression ratio of the engine.
22. **horsepower:** Horsepower of the car.
23. **peakrpm:** Peak revolutions per minute.
24. **citympg:** Miles per gallon in the city.
25. **highwaympg:** Miles per gallon on the highway.
26. **price:** Price of the car.

6.2.2 Check Unique Values for each variable.

```
[122]: # Check Unique Values for each variable.
for column in cars_df.columns:
    unique_values = cars_df[column].unique()
    print(f"Unique values in {column}:\n{unique_values}\n")
```

Unique values in car_ID:

```
[ 1  2  3  4  5  6  7  8  9 10 11 12 13 14 15 16 17 18
 19 20 21 22 23 24 25 26 27 28 29 30 31 32 33 34 35 36
 37 38 39 40 41 42 43 44 45 46 47 48 49 50 51 52 53 54
 55 56 57 58 59 60 61 62 63 64 65 66 67 68 69 70 71 72
 73 74 75 76 77 78 79 80 81 82 83 84 85 86 87 88 89 90
 91 92 93 94 95 96 97 98 99 100 101 102 103 104 105 106 107 108
109 110 111 112 113 114 115 116 117 118 119 120 121 122 123 124 125 126
127 128 129 130 131 132 133 134 135 136 137 138 139 140 141 142 143 144
145 146 147 148 149 150 151 152 153 154 155 156 157 158 159 160 161 162
163 164 165 166 167 168 169 170 171 172 173 174 175 176 177 178 179 180
181 182 183 184 185 186 187 188 189 190 191 192 193 194 195 196 197 198
199 200 201 202 203 204 205]
```

Unique values in symboling:

```
[ 3  1  2  0 -1 -2]
```

Unique values in CarName:

```

['alfa-romero giulia' 'alfa-romero stelvio' 'alfa-romero Quadrifoglio'
'audi 100 ls' 'audi 100ls' 'audi fox' 'audi 5000' 'audi 4000'
'audi 5000s (diesel)' 'bmw 320i' 'bmw x1' 'bmw x3' 'bmw z4' 'bmw x4'
'bmw x5' 'chevrolet impala' 'chevrolet monte carlo' 'chevrolet vega 2300'
'dodge rampage' 'dodge challenger se' 'dodge d200' 'dodge monaco (sw)'
'dodge colt hardtop' 'dodge colt (sw)' 'dodge coronet custom'
'dodge dart custom' 'dodge coronet custom (sw)' 'honda civic'
'honda civic cvcc' 'honda accord cvcc' 'honda accord lx'
'honda civic 1500 gl' 'honda accord' 'honda civic 1300' 'honda prelude'
'honda civic (auto)' 'isuzu MU-X' 'isuzu D-Max ' 'isuzu D-Max V-Cross'
'jaguar xj' 'jaguar xf' 'jaguar xk' 'maxda rx3' 'maxda glc deluxe'
'mazda rx2 coupe' 'mazda rx-4' 'mazda glc deluxe' 'mazda 626' 'mazda glc'
'mazda rx-7 gs' 'mazda glc 4' 'mazda glc custom l' 'mazda glc custom'
'buick electra 225 custom' 'buick century luxus (sw)' 'buick century'
'buick skyhawk' 'buick opel isuzu deluxe' 'buick skylark'
'buick century special' 'buick regal sport coupe (turbo)'
'mercury cougar' 'mitsubishi mirage' 'mitsubishi lancer'
'mitsubishi outlander' 'mitsubishi g4' 'mitsubishi mirage g4'
'mitsubishi montero' 'mitsubishi pajero' 'Nissan versa' 'nissan gt-r'
'nissan rogue' 'nissan latio' 'nissan titan' 'nissan leaf' 'nissan juke'
'nissan note' 'nissan clipper' 'nissan nv200' 'nissan dayz' 'nissan fuga'
'nissan otti' 'nissan teana' 'nissan kicks' 'peugeot 504' 'peugeot 304'
'peugeot 504 (sw)' 'peugeot 604sl' 'peugeot 505s turbo diesel'
'plymouth fury iii' 'plymouth cricket' 'plymouth satellite custom (sw)'
'plymouth fury gran sedan' 'plymouth valiant' 'plymouth duster'
'porsche macan' 'porsche panamera' 'porsche cayenne' 'porsche boxer'
'renault 12tl' 'renault 5 gtl' 'saab 99e' 'saab 99le' 'saab 99gle'
'subaru' 'subaru dl' 'subaru brz' 'subaru baja' 'subaru r1' 'subaru r2'
'subaru trezia' 'subaru tribeca' 'toyota corona mark ii' 'toyota corona'
'toyota corolla 1200' 'toyota corona hardtop' 'toyota corolla 1600 (sw)'
'toyota carina' 'toyota mark ii' 'toyota corolla'
'toyota corolla liftback' 'toyota celica gt liftback'
'toyota corolla tercel' 'toyota corona liftback' 'toyota starlet'
'toyota tercel' 'toyota cressida' 'toyota celica gt' 'toyota tercel'
'volkswagen rabbit' 'volkswagen 113i deluxe sedan' 'volkswagen model 111'
'volkswagen type 3' 'volkswagen 411 (sw)' 'volkswagen super beetle'
'volkswagen dasher' 'vw dasher' 'vw rabbit' 'volkswagen rabbit'
'volkswagen rabbit custom' 'volvo 145e (sw)' 'volvo 144ea' 'volvo 244dl'
'volvo 245' 'volvo 264gl' 'volvo diesel' 'volvo 246']

```

Unique values in fueltype:

```
['gas' 'diesel']
```

Unique values in aspiration:

```
['std' 'turbo']
```

Unique values in doornumber:

```
['two' 'four']
```

Unique values in carbody:

['convertible' 'hatchback' 'sedan' 'wagon' 'hardtop']

Unique values in drivewheel:

['rwd' 'fwd' '4wd']

Unique values in enginelocation:

['front' 'rear']

Unique values in wheelbase:

[88.6 94.5 99.8 99.4 105.8 99.5 101.2 103.5 110. 88.4 93.7 103.3
 95.9 86.6 96.5 94.3 96. 113. 102. 93.1 95.3 98.8 104.9 106.7
115.6 96.6 120.9 112. 102.7 93. 96.3 95.1 97.2 100.4 91.3 99.2
107.9 114.2 108. 89.5 98.4 96.1 99.1 93.3 97. 96.9 95.7 102.4
102.9 104.5 97.3 104.3 109.1]

Unique values in carlength:

[168.8 171.2 176.6 177.3 192.7 178.2 176.8 189. 193.8 197. 141.1 155.9
158.8 157.3 174.6 173.2 144.6 150. 163.4 157.1 167.5 175.4 169.1 170.7
172.6 199.6 191.7 159.1 166.8 169. 177.8 175. 190.9 187.5 202.6 180.3
208.1 199.2 178.4 173. 172.4 165.3 170.2 165.6 162.4 173.4 181.7 184.6
178.5 186.7 198.9 167.3 168.9 175.7 181.5 186.6 156.9 157.9 172. 173.5
173.6 158.7 169.7 166.3 168.7 176.2 175.6 183.5 187.8 171.7 159.3 165.7
180.2 183.1 188.8]

Unique values in carwidth:

[64.1 65.5 66.2 66.4 66.3 71.4 67.9 64.8 66.9 70.9 60.3 63.6 63.8 64.6
63.9 64. 65.2 62.5 66. 61.8 69.6 70.6 64.2 65.7 66.5 66.1 70.3 71.7
70.5 72. 68. 64.4 65.4 68.4 68.3 65. 72.3 66.6 63.4 65.6 67.7 67.2
68.9 68.8]

Unique values in carheight:

[48.8 52.4 54.3 53.1 55.7 55.9 52. 53.7 56.3 53.2 50.8 50.6 59.8 50.2
52.6 54.5 58.3 53.3 54.1 51. 53.5 51.4 52.8 47.8 49.6 55.5 54.4 56.5
58.7 54.9 56.7 55.4 54.8 49.4 51.6 54.7 55.1 56.1 49.7 56. 50.5 55.2
52.5 53. 59.1 53.9 55.6 56.2 57.5]

Unique values in curbweight:

[2548 2823 2337 2824 2507 2844 2954 3086 3053 2395 2710 2765 3055 3230
3380 3505 1488 1874 1909 1876 2128 1967 1989 2191 2535 2811 1713 1819
1837 1940 1956 2010 2024 2236 2289 2304 2372 2465 2293 2734 4066 3950
1890 1900 1905 1945 1950 2380 2385 2500 2410 2443 2425 2670 2700 3515
3750 3495 3770 3740 3685 3900 3715 2910 1918 1944 2004 2145 2370 2328
2833 2921 2926 2365 2405 2403 1889 2017 1938 1951 2028 1971 2037 2008
2324 2302 3095 3296 3060 3071 3139 3020 3197 3430 3075 3252 3285 3485
3130 2818 2778 2756 2800 3366 2579 2460 2658 2695 2707 2758 2808 2847
2050 2120 2240 2190 2340 2510 2290 2455 2420 2650 1985 2040 2015 2280]

```
3110 2081 2109 2275 2094 2122 2140 2169 2204 2265 2300 2540 2536 2551
2679 2714 2975 2326 2480 2414 2458 2976 3016 3131 3151 2261 2209 2264
2212 2319 2254 2221 2661 2563 2912 3034 2935 3042 3045 3157 2952 3049
3012 3217 3062]
```

Unique values in enginetype:

```
['dohc' 'ohcv' 'ohc' 'l' 'rotor' 'ohcf' 'dohcv']
```

Unique values in cylindernumber:

```
['four' 'six' 'five' 'three' 'twelve' 'two' 'eight']
```

Unique values in enginesize:

```
[130 152 109 136 131 108 164 209 61 90 98 122 156 92 79 110 111 119
258 326 91 70 80 140 134 183 234 308 304 97 103 120 181 151 194 203
132 121 146 171 161 141 173 145]
```

Unique values in fuelsystem:

```
['mpfi' '2bbl' 'mfi' '1bbl' 'spfi' '4bbl' 'idi' 'spdi']
```

Unique values in boreratio:

```
[3.47 2.68 3.19 3.13 3.5 3.31 3.62 2.91 3.03 2.97 3.34 3.6 2.92 3.15
3.43 3.63 3.54 3.08 3.33 3.39 3.76 3.58 3.46 3.8 3.78 3.17 3.35 3.59
2.99 3.7 3.61 3.94 3.74 2.54 3.05 3.27 3.24 3.01]
```

Unique values in stroke:

```
[2.68 3.47 3.4 2.8 3.19 3.39 3.03 3.11 3.23 3.46 3.9 3.41
3.07 3.58 4.17 2.76 3.15 3.255 3.16 3.64 3.1 3.35 3.12 3.86
3.29 3.27 3.52 2.19 3.21 2.9 2.07 2.36 2.64 3.08 3.5 3.54
2.87 ]
```

Unique values in compressionratio:

```
[ 9. 10. 8. 8.5 8.3 7. 8.8 9.5 9.6 9.41 9.4 7.6
9.2 10.1 9.1 8.1 11.5 8.6 22.7 22. 21.5 7.5 21.9 7.8
8.4 21. 8.7 9.31 9.3 7.7 22.5 23. ]
```

Unique values in horsepower:

```
[111 154 102 115 110 140 160 101 121 182 48 70 68 88 145 58 76 60
86 100 78 90 176 262 135 84 64 120 72 123 155 184 175 116 69 55
97 152 200 95 142 143 207 288 73 82 94 62 56 112 92 161 156 52
85 114 162 134 106]
```

Unique values in peakrpm:

```
[5000 5500 5800 4250 5400 5100 4800 6000 4750 4650 4200 4350 4500 5200
4150 5600 5900 5750 5250 4900 4400 6600 5300]
```

Unique values in citympg:

```
[21 19 24 18 17 16 23 20 15 47 38 37 31 49 30 27 25 13 26 36 22 14 45 28
32 35 34 29 33]
```

Unique values in highwaympg:

```
[27 26 30 22 25 20 29 28 53 43 41 38 24 54 42 34 33 31 19 17 23 32 39 18
 16 37 50 36 47 46]
```

Unique values in price:

```
[13495.    16500.    13950.    17450.    15250.    17710.    18920.
 23875.    17859.167 16430.    16925.    20970.    21105.    24565.
 30760.    41315.    36880.    5151.     6295.     6575.     5572.
  6377.     7957.     6229.    6692.     7609.     8558.     8921.
12964.     6479.     6855.     5399.     6529.     7129.     7295.
  7895.     9095.     8845.    10295.    12945.    10345.     6785.
 8916.5    11048.    32250.    35550.    36000.     5195.     6095.
  6795.     6695.     7395.    10945.    11845.    13645.    15645.
  8495.    10595.    10245.    10795.    11245.    18280.    18344.
25552.    28248.    28176.    31600.    34184.    35056.    40960.
45400.    16503.     5389.     6189.     6669.     7689.     9959.
  8499.    12629.    14869.    14489.     6989.     8189.     9279.
  5499.     7099.     6649.     6849.     7349.     7299.     7799.
  7499.     7999.     8249.     8949.     9549.    13499.    14399.
17199.    19699.    18399.    11900.    13200.    12440.    13860.
15580.    16900.    16695.    17075.    16630.    17950.    18150.
12764.    22018.    32528.    34028.    37028.    31400.5    9295.
  9895.    11850.    12170.    15040.    15510.    18620.     5118.
  7053.     7603.     7126.     7775.     9960.     9233.    11259.
  7463.    10198.     8013.    11694.     5348.     6338.     6488.
  6918.     7898.     8778.     6938.     7198.     7788.     7738.
  8358.     9258.     8058.     8238.     9298.     9538.     8449.
  9639.     9989.    11199.    11549.    17669.     8948.    10698.
  9988.    10898.    11248.    16558.    15998.    15690.    15750.
  7975.     7995.     8195.     9495.     9995.    11595.     9980.
13295.    13845.    12290.    12940.    13415.    15985.    16515.
18420.    18950.    16845.    19045.    21485.    22470.    22625.    ]
```

6.3 3. Data Vizualization, Storytelling & Experimenting with charts : Understand the relationships between variables

Chart - 1 Histogram (Checking Distribution of dataset Independent Variables)

```
[123]: numeric_columns = cars_df.select_dtypes(include=['number'])

# Drop 'car_ID' and 'symboling' columns
columns_to_plot = numeric_columns.drop(['car_ID', 'price'], axis=1)

# Create a histogram with KDE and lines for mean, median, and mode
plt.figure(figsize=(15, 10))
```

```

for i, column in enumerate(columns_to_plot.columns):
    plt.subplot(4, 4, i+1)
    sns.histplot(cars_df[column], kde=True, bins=20, alpha=0.7, color = 'blue')

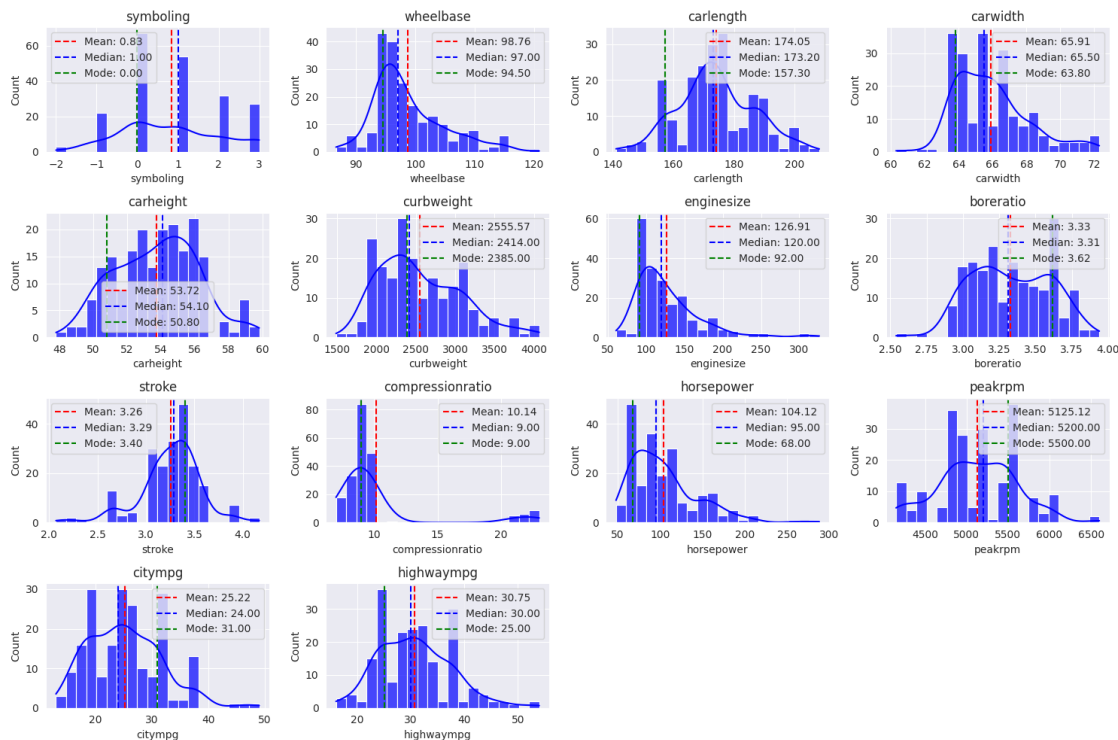
    mean_val = cars_df[column].mean()
    median_val = cars_df[column].median()
    mode_val = cars_df[column].mode()[0]

    plt.axvline(mean_val, color='red', linestyle='--', label=f'Mean: {mean_val:.2f}')
    plt.axvline(median_val, color='blue', linestyle='--', label=f'Median: {median_val:.2f}')
    plt.axvline(mode_val, color='green', linestyle='--', label=f'Mode: {mode_val:.2f}')

    plt.title(column)
    plt.legend()

plt.tight_layout()
plt.show()

```



[124]: cars_df.skew()

<ipython-input-124-6bac9c360b6d>:1: FutureWarning: The default value of numeric_only in DataFrame.skew is deprecated. In a future version, it will default to False. In addition, specifying 'numeric_only=None' is deprecated. Select only valid columns or specify the value of numeric_only to silence this warning.

```
cars_df.skew()
```

```
[124]: car_ID          0.000000
      symboling       0.211072
      wheelbase      1.050214
      carlength      0.155954
      carwidth       0.904003
      carheight      0.063123
      curbweight     0.681398
      enginesize      1.947655
      boreratio      0.020156
      stroke        -0.689705
      compressionratio 2.610862
      horsepower     1.405310
      peakrpm        0.075159
      citympg        0.663704
      highwaympg     0.539997
      price          1.777678
      dtype: float64
```

1. Why did you pick the specific chart? This specific chart was picked up to check the if the numerical columns which are to be involved in the building the **Regression** model follows **Gaussian Distribution** or not.

In general, in a particular dataset

Mean < Median < Mode - if the dataset follows **Left Skewed Distribution**

Mean > Median < Mode - if the dataset follows **Right Skewed Distribution**

Mean = Median = Mode - if the dataset follows **Symmetric Distribution**

2. What is/are the insight(s) found from the chart? The dataset exhibits a mix of symmetric and skewed distributions:

1. **Symmetric Distributions:** car_ID, carlength, carheight, boreratio, peakrpm, and stroke have skewness values close to zero, indicating relatively symmetric distributions.
2. **Right-Skewed Distributions:** Variables such as wheelbase, carwidth, curbweight, enginesize, compressionratio, horsepower, citympg, highwaympg, and price display right-skewed distributions, suggesting longer tails on the right.
3. **Left-Skewed Distribution:** The variable stroke has a negative skewness value, indicating a left-skewed distribution with a longer left tail.

The above skewness in data needs to be treated because for specific machine learning algorithms such as linear regression works well if the data is symmetric i.e normally distributed. Hence for better model performance skewness should be reduced and we will see that in further stage of the project.

[124]:

Chart - 2 Countplot (Checking Distribution of dataset Independent Variables)

```
[125]: # Exclude 'CarName' from categorical variables
cat_vars = [var for var in cars_df.select_dtypes(include='object').columns if
    var != 'CarName']

# Determine the number of rows and columns for subplots
num_rows = (len(cat_vars) + 1) // 2 # Ensuring an extra row if the number of
    variables is odd
num_cols = 2 # Adjust this based on your preference


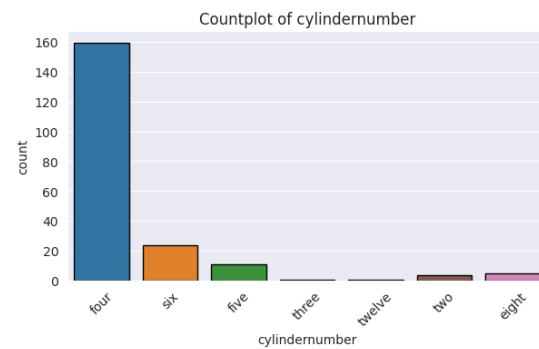
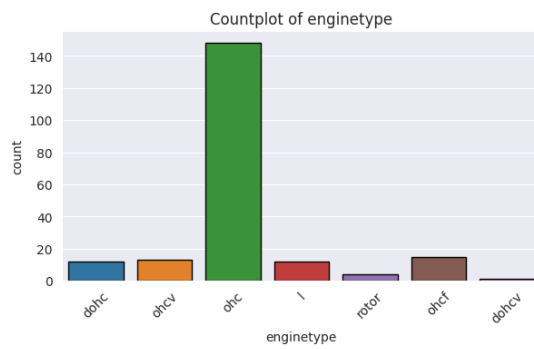
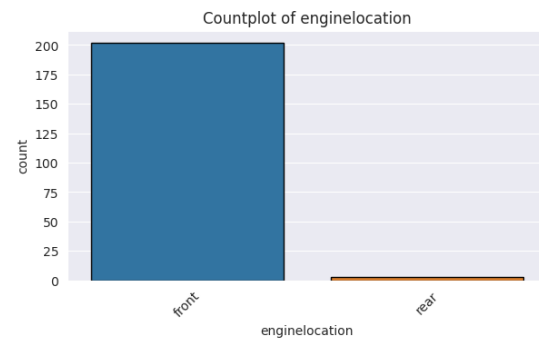
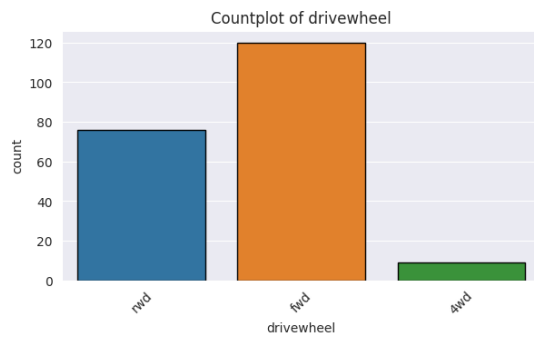
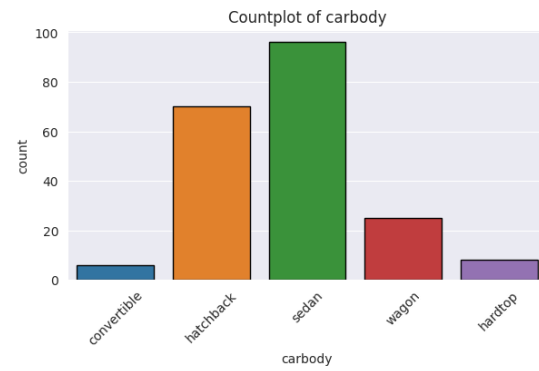
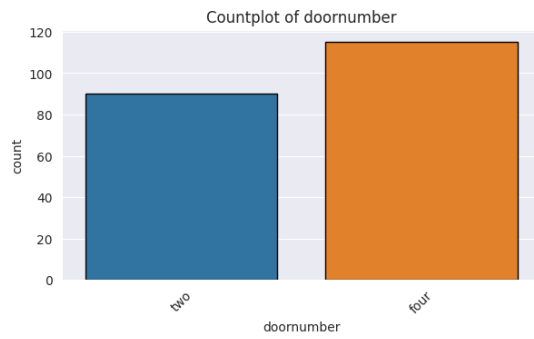
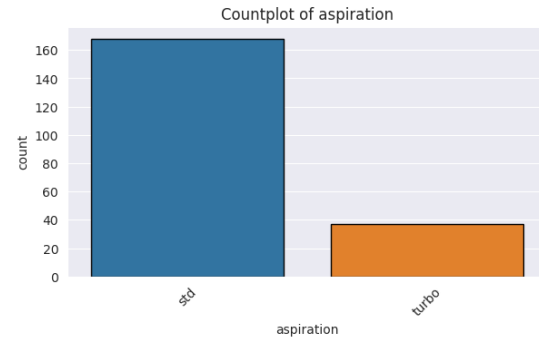
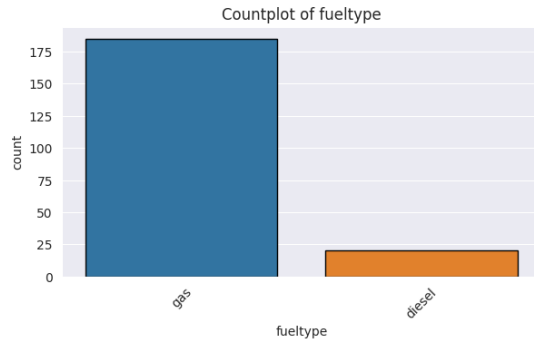
# Set up subplots
fig, axes = plt.subplots(num_rows, num_cols, figsize=(15, 5 * num_rows))
fig.subplots_adjust(hspace=0.5) # Adjust vertical spacing

# Flatten the axes array to simplify indexing
axes = axes.flatten()

# Loop through each categorical variable (excluding 'CarName') and create a
    countplot
for i, var in enumerate(cat_vars):
    sns.countplot(x=var, data=cars_df, ax=axes[i], edgecolor='black')
    axes[i].set_title(f'Countplot of {var}')
    axes[i].tick_params(axis='x', rotation=45) # Rotate x-axis labels for
        better readability

# If the number of subplots is odd, remove the empty subplot
if len(cat_vars) % 2 != 0:
    fig.delaxes(axes[-1])

# Display the subplots
plt.show()
```



1. Why did you pick the specific chart? The countplot is a preferred choice for visualizing categorical data due to its simplicity and effectiveness. Designed specifically for categorical variables, it provides a clear representation of the frequency of each category through bar heights.

2. What is/are the insight(s) found from the chart? The data reveals that petrol is the most prevalent fueltype, followed by diesel, CNG, and hybrid. Natural aspiration dominates the aspiration category, followed by turbo. The majority of vehicles have four doors, followed by five and two. Hatchbacks are the most common carbody, followed by sedans and SUVs. Four-wheel drive is the most common drivewheel type, followed by rear-wheel drive and front-wheel drive. Front-engine placement is more common than rear-engine placement. The majority of engines are four-cylinder, followed by three-cylinder, six-cylinder, and five-cylinder engines. The most common number of cylinders is four, followed by three and five. MPFI is the most common fuel system, followed by SPFI and EFI. These insights provide a comprehensive understanding of the data's characteristics and enable predictions about future trends.

Chart - 3 Histogram (Checking Distribution of dataset Dependent Variable)

```
[126]: # Set up the figure size
plt.figure(figsize=(10, 6))

# Plotting the histogram for 'price' variable
sns.histplot(cars_df['price'], bins=30, kde=True, color='blue',
             ↪edgecolor='black')

# Plotting the mean line
plt.axvline(cars_df['price'].mean(), color='red', linestyle='dashed',
             ↪linewidth=2, label='Mean')


# Plotting the median line
plt.axvline(cars_df['price'].median(), color='green', linestyle='dashed',
             ↪linewidth=2, label='Median')

# Plotting the mode line
plt.axvline(cars_df['price'].mode().values[0], color='orange',
             ↪linestyle='dashed', linewidth=2, label='Mode')

# Adding labels and title
plt.xlabel('Price')
plt.ylabel('Frequency')
plt.title('Histogram of Price with Mean, Median, and Mode')

# Adding a legend
plt.legend()
```

```
# Displaying the plot
plt.show()
```



1. Why did you pick the specific chart? This specific chart was picked up to check the if the numerical columns which are to be involved in the building the **Regression** model follows the **Gaussian Distribution** or not.

In general, in a particular dataset

Mean < Median < Mode - if the dataset follows **Left Skewed Distribution**

Mean > Median < Mode - if the dataset follows **Right Skewed Distribution**

Mean = Median = Mode - if the dataset follows **Symmetric Distribution**

2. What is/are the insight(s) found from the chart? When we talk particularly about the target variable i.e 'price' even this has a skewed distribution i.e **right skewed** distribution. This needs to be treated as we know that a skewed dataset especially in case of regression models such as linear regression algorithms lowers the models performance.

Chart - 4 Boxplot (Identifying Outliers for dataset) Implemented a robust outlier detection technique, leveraging the Interquartile Range (IQR) method. This method ensures the identification of unexpected extreme values in the dataset tails.

Percentile Computation:

Calculated the first quartile (Q1, 25th percentile) and the third quartile (Q3, 75th percentile) to establish the data spread.

Bound Definition:

Defined lower and upper bounds using Q1, Q3, and the Interquartile Range (IQR):

Lower Bound: $Q1 - 1.5 \times IQR$

Upper Bound: $Q3 + 1.5 \times IQR$

Outlier Identification: Identified outliers as data points falling below the lower bound or above the upper bound.

```
[127]: # Count the number of outliers for each feature variable
outliers_count = {}

# Iterate through each column in the dataframe
for column in cars_df.columns:
    # Check if the data type of the column is numeric (integer or float)
    if cars_df[column].dtype in ['int64', 'float64']:
        # Calculate the first quartile (Q1)
        Q1 = cars_df[column].quantile(0.25)

        # Calculate the third quartile (Q3)
        Q3 = cars_df[column].quantile(0.75)

        # Calculate the Interquartile Range (IQR)
        IQR = Q3 - Q1

        # Calculate the lower and upper bounds to identify outliers
        lower_bound = Q1 - 1.5 * IQR
        upper_bound = Q3 + 1.5 * IQR

        # Count the number of outliers for the current column
        outliers_count[column] = cars_df[(cars_df[column] < lower_bound) |
↪(cars_df[column] > upper_bound)].shape[0]

# Display the count of outliers for each feature variable
print("Number of outliers for each feature variable:")
outliers_count
```

Number of outliers for each feature variable:

```
[127]: {'car_ID': 0,
        'symboling': 0,
        'wheelbase': 3,
        'carlength': 1,
        'carwidth': 8,
        'carheight': 0,
```

```

'curbweight': 0,
'engine size': 10,
'bore ratio': 0,
'stroke': 20,
'compressionratio': 28,
'horsepower': 6,
'peakrpm': 2,
'citympg': 2,
'highwaympg': 3,
'price': 15}

```

```

[128]: # Extracting numerical variables
numeric_variables = cars_df.select_dtypes(include=['float64', 'int64']).columns

# Removing 'price' column from the list of numerical variables
numeric_variables = numeric_variables.drop('price')















# Setting up subplots
fig, axes = plt.subplots(nrows=len(numeric_variables), ncols=1, figsize=(7, 5 *
    len(numeric_variables)))
fig.subplots_adjust(hspace=0.5)

# Plotting box plots for each numeric variable
for i, variable in enumerate(numeric_variables):
    # Creating a box plot for the current numeric variable
    sns.boxplot(data=cars_df, x=variable, ax=axes[i], color = 'blue')

    # Setting the title of the subplot
    axes[i].set_title(f'Boxplot of {variable}')

# Displaying the subplots
plt.show()

```



```
[129]: for column in cars_df.columns:
        if cars_df[column].dtype in ['int64', 'float64']:
            Q1 = cars_df[column].quantile(0.25)
            Q3 = cars_df[column].quantile(0.75)
            IQR = Q3 - Q1

            lower_bound = Q1 - 1.5 * IQR
            upper_bound = Q3 + 1.5 * IQR

            cars_df.loc[cars_df[column] < lower_bound, column] = lower_bound
            cars_df.loc[cars_df[column] > upper_bound, column] = upper_bound
```

1. Why did you pick the specific chart? The selection of this specific chart aimed to visually assess outliers within the independent numeric variables. A boxplot was utilized as the chosen visualization method for this examination.

2. What is/are the insight(s) found from the chart? From the visual analysis, it is evident that certain feature variables exhibit outliers, identified using the IQR approach. These outliers extend beyond the upper and lower bounds of the expected range. Subsequently, a robust data treatment strategy was implemented, replacing the identified outliers with their respective upper and lower bound values. This intervention ensures data integrity and aligns the feature variables with a more standardized and reliable distribution.

Chart - 5 Regplot (For Assessing the Correlation Between dependent and independent Variables)

```
[130]: # Calculate the correlation matrix for all numeric variables with respect to
        ↪ 'price'
        price_correlation = cars_df.corr()['price']

        # Displaying the correlation matrix
        price_correlation
```

<ipython-input-130-2787b5888632>:2: FutureWarning: The default value of numeric_only in DataFrame.corr is deprecated. In a future version, it will default to False. Select only valid columns or specify the value of numeric_only to silence this warning.

```
price_correlation = cars_df.corr()['price']
```

```
[130]: car_ID          -0.089603
        symboling      -0.092705
        wheelbase       0.595909
        carlength       0.712455
        carwidth        0.783230
        carheight       0.142033
```



```

curbweight      0.864597
engine_size     0.860063
bore_ratio      0.572685
stroke          0.073830
compression_ratio -0.056573
horsepower      0.821715
peak_rpm        -0.088630
city_mpg        -0.718290
highway_mpg     -0.733692
price           1.000000
Name: price, dtype: float64

```

```

[131]: # Extracting numerical variables
numeric_variables = cars_df.select_dtypes(include=['float64', 'int64']).columns

# Removing 'car_ID' and 'symboling' columns from the list of numerical variables
numeric_variables = numeric_variables.drop(['car_ID', 'price'])

# Defining dependent variable
dependent_variable = 'price'

# Setting up subplots
total_plots = len(numeric_variables)
num_cols = 3
num_rows = (total_plots - 1) // num_cols + 1
plt.figure(figsize=(15, 5 * num_rows))

# Plotting regression plots for each numeric variable
for i, column in enumerate(numeric_variables):
    plt.subplot(num_rows, num_cols, i+1)

    # Creating a regression plot for the current numeric variable against the
    # dependent variable
    sns.regplot(x=cars_df[column], y=cars_df[dependent_variable],
                scatter_kws={"color": 'orange'}, line_kws={"color": "black"})

    # Setting the title of the subplot
    plt.title(f'{column} vs. {dependent_variable}')

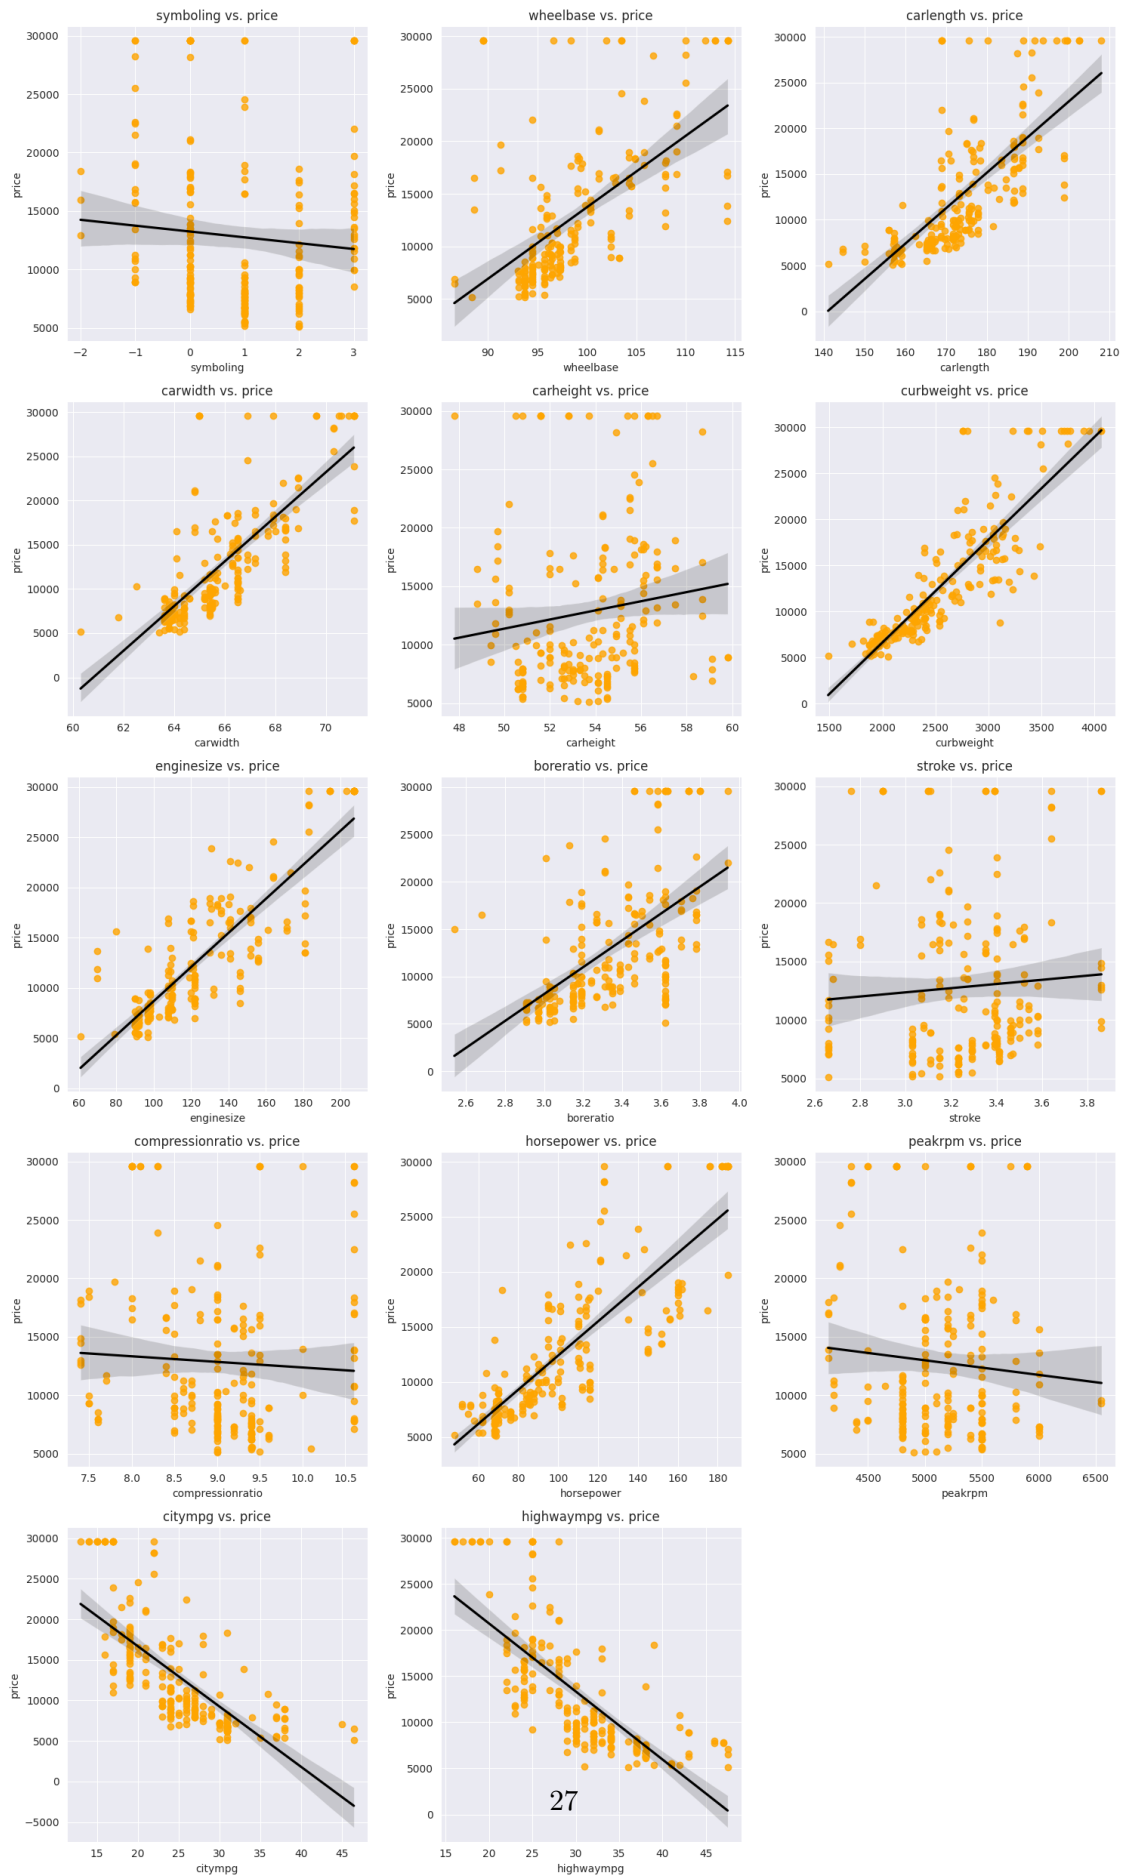
    # Setting x-axis label
    plt.xlabel(column)

    # Setting y-axis label
    plt.ylabel(dependent_variable)

# Adjusting layout for better visualization
plt.tight_layout()

```

```
# Displaying the subplots  
plt.show()
```



1. Why did you pick the specific chart? The application of **regplot** is fitting for our visualization needs as it not only facilitates the exploration of the linear relationship between independent and dependent variables but also incorporates a trendline. This feature enhances our understanding by visually representing the best-fit linear regression line, elucidating how variations in the independent variables correspond to changes in the dependent variable.

2. What is/are the insight(s) found from the chart? The following insights were found from the above visualization :

1. Strong Positive Correlation:

For variables like ‘carwidth’, ‘curbweight’, ‘enginesize’, and ‘horsepower’ that exhibit strong positive correlations with ‘price’, the regplot with a trendline will show a clear upward-sloping line. As these variables increase, the ‘price’ tends to increase, forming a distinct positive linear relationship captured by the trendline.

2. Moderate Positive Correlation:

Variables with moderate positive correlations such as ‘wheelbase’, ‘carlength’, and ‘boreratio’ will also result in a positive slope in the regplot with a trendline. The slope may not be as steep as in the case of strong positive correlations, indicating a moderately positive linear relationship represented by the trendline.

3. Weak Positive Correlation:

For variables like ‘carheight’, ‘compressionratio’, and ‘peakrpm’ with weak positive correlations, the regplot with a trendline may show a positive slope, but the relationship is not as pronounced. The points on the scatter plot may not form a clear linear trend, and the trendline captures the subtle positive relationship.

4. Negative Correlation:

Variables ‘citympg’ and ‘highwaympg’ with negative correlations will yield a regplot with a downward-sloping trendline. As these variables increase, the ‘price’ tends to decrease, showcasing a negative linear relationship captured by the trendline.

5. Weak Negative Correlation:

Variables ‘symboling’ and ‘stroke’ exhibit weak negative correlations. The regplot with a trendline may show a negative slope, but the relationship is not strongly evident. The trendline captures the subtle negative relationship, and the scatter of points may not form a distinct downward trend.

Chart - 6 Correlation HeatMap (Assessing Correlation amongst all the variables)

```
[132]: # Chart - 4 visualization code
## Correlation

# Setting up the figure size
plt.figure(figsize=(15, 8))


# Calculating the correlation matrix
correlation = cars_df.corr()

# Creating a heatmap to visualize the correlation matrix
sns.heatmap(correlation, annot=True, cmap='coolwarm')

# Displaying the plot
plt.show()
```

<ipython-input-132-6b03ad9b8594>:8: FutureWarning: The default value of numeric_only in DataFrame.corr is deprecated. In a future version, it will default to False. Select only valid columns or specify the value of numeric_only to silence this warning.

```
correlation = cars_df.corr()
```



1. Why did you pick the specific chart? A correlation heatmap is used to visualize the strength and direction of relationships between two or more variables in a dataset. The heatmap

displays correlation coefficients as color-coded values, allowing for easy identification of patterns and insights into the associations among variables.

2. What is/are the insight(s) found from the chart? The correlation matrix effectively visualizes the strength and direction of correlations between variables. The color scale employed in the matrix enhances this understanding, where dark red indicates a strong positive correlation, and dark blue represents a substantial negative correlation. The intensity of the color corresponds to the strength of the correlation, with darker shades signifying stronger associations. As the color transitions to lighter shades in either red or blue, the strength of the correlation diminishes. This visual representation provides a quick and insightful overview of the relationships within the dataset.

Focusing on the correlations between some notable variables:

Horsepower (hp) and curb weight: A strong positive correlation is observed, suggesting that as horsepower increases, so does curb weight. This is intuitive, as heavier vehicles typically require more powerful engines.

Horsepower (hp) and city mpg: A moderate negative correlation is seen, implying that as horsepower increases, city fuel efficiency tends to decrease. This is expected, as more powerful engines generally consume more fuel.

Wheelbase and car length: A strong positive correlation is evident, indicating that as wheelbase increases, car length also tends to increase. This is because wheelbase is the distance between the front and rear axles, and a larger wheelbase typically translates to a longer car.

Bore ratio and stroke: A moderate negative correlation is evident, suggesting that as bore ratio increases (meaning the cylinder is wider relative to its stroke), stroke tends to decrease (meaning the piston travels a shorter distance). This relationship is often observed in engine design, as bore and stroke are crucial factors in determining engine characteristics.

Compression ratio and horsepower (hp): A moderate positive correlation is seen, indicating that as compression ratio increases, horsepower also tends to increase. This is because a higher compression ratio allows for more efficient combustion, leading to higher power output.

6.4 4. Feature Engineering & Data Pre-processing

6.4.1 1. Handling Missing Values

[133]: `# Handling Missing Values & Missing Value Imputation`

What all missing value imputation techniques have you used and why did you use those techniques? The dataset utilized for this project, namely `cars_df`, demonstrates a noteworthy attribute—complete absence of missing values. Given this, there is no need for imputation or handling of **missing data**, allowing us to seamlessly progress to the next step, which involves treating outliers in the dataset.

6.4.2 2. Handling Outliers

Having successfully identified outliers within the feature variables using the **IQR** approach, the ensuing step involves implementing outlier treatment for the following list of variables.

1. wheelbase
2. carlength
3. carwidth
4. enginesize
5. stroke
6. compressionratio
7. horsepower
8. peakrpm
9. citympg
10. highwaympg
11. price

```
[134]: for column in cars_df.columns:
        if cars_df[column].dtype in ['int64', 'float64']:
            Q1 = cars_df[column].quantile(0.25)
            Q3 = cars_df[column].quantile(0.75)
            IQR = Q3 - Q1

            lower_bound = Q1 - 1.5 * IQR
            upper_bound = Q3 + 1.5 * IQR

            cars_df.loc[cars_df[column] < lower_bound, column] = lower_bound
            cars_df.loc[cars_df[column] > upper_bound, column] = upper_bound
```

What all outlier treatment techniques have you used and why did you use those techniques? Percentile Computation:

Calculate Q1 (25th percentile) and Q3 (75th percentile) of the dataset. Bound Definition:

Establish lower and upper bounds using the IQR:

Lower Bound: $Q1 - 1.5 \times IQR$

Upper Bound: $Q3 + 1.5 \times IQR$

Outlier Replacement:

Replace values that fall outside the upper or lower bounds with the respective boundary values.

The IQR methodology, being robust and distribution-free, furnishes a dependable means of identifying outliers. Temporary removal of these outliers unveils more stable structures and relationships in the data that may be obscured by extreme values. Importantly, no data is permanently lost, as outliers are reintroduced into the dataset post-analysis. This approach circumvents assumptions about the underlying distribution, allowing the analysis to unveil central tendencies that outliers may otherwise conceal.

6.4.3 3. Feature Manipulation

After having a look at the dataset, We can add three more features(variables) from the existing dataset. 1. **mileage** = (0.6 * city miles per gallon) + (0.4 * highway miles per gallon)

2. **car_area** = carlength * carwidth

3. **car_volume** = carlength * carwidth * carheight

```
[135]: # Manipulate Features to minimize feature correlation and create new features
cars_df['mileage'] = 0.6*cars_df['citympg'] + 0.4*cars_df['highwaympg']
cars_df['car_area'] = cars_df['carlength']*cars_df['carwidth']
cars_df['car_volume'] = cars_df['carlength'] * cars_df['carwidth'] *
↳cars_df['carheight']
```

The dataset also has **CarName** column which contains the name of the car along with the company name. From this we can extract the company name in a separate column.

```
[136]: cars_df['company'] = cars_df['CarName'].str.split(" ", expand=True)[0]
```

Also when you pay closer attention to the dataset, you will find that some of the company names have been wrongly entered in the dataset with incorrect spellings due to which they are being considered as separate company which is not the case. Therefore we will be doing some modification in that by correcting the names of the company.

```
[137]: cars_df['company'] = cars_df['company'].replace({'toyouta': 'Toyota', 'vw':
↳'Volkswagen', 'vokswagen': 'Volkswagen', 'maxda': 'Mazda', 'porcshce': 'Porsche'})
cars_df['company'] = cars_df['company'].str.title()
```

Removing the irrelevant columns from the context of Model Implementation

From the above dataset we will be removing two columns i.e Car_ID & CarName. The reason for dropping these two columns from the dataset is their relevance in building and implementing the machine learning model. Firstly talking about Car_ID though being a numeric column, it is just representing a particular id assigned to a particular car and there's no actual relationship between that one could assess between it and price of a particular car.

Secondly talking about the CarName it is a categorical variable which contains the names of the different cars. Here we have used this column indirectly by extracting the company name from this column and then later converting that to a Frequency Encoded column. Also car name directly cannot be used for building and implementing regression models as it contains text values. Also we will not be dropping the **CarName** column in this step because of we have some (calculations/operations to perform before we actually drop it).

1. **Car_ID**

2. **CarName**

```
[138]: # Dropping the irrelevant columns
cars_df.drop(['car_ID'], axis = 1, inplace = True)
```


6.4.4 4. Categorical Encoding

List of Variables which will have **Frequency Encoding** :-

1. company

List of Variables which will have **Label Encoding** :-

1. doornumber
2. cylindernumber

List of Variables which will have **One Hot Encoding** :-

1. carbody
2. fueltype
3. aspiration
4. enginelocation
5. enginetype
6. fuelsystem
7. drivewheel

What all categorical encoding techniques have you used & why did you use those techniques? In the above case I have used (**Frequency Encoding**), (**Label Encoding**) & (**One Hot Encoding**) as the variables were of such nature that the above three were most relevant options. For variable such as **carbody**, **fueltype**, **aspiration** etc. since they are **nominal** data that have no inherent **order/ranking** attached to them, so i went ahead One Hot Encoding. For variable such as **doornumber**, **cylindernumber** etc. Label Encoding since they were variables which can be represented in two numeric values. For the variable **company** name I have used the Frequency Encoding because it contains the 31 distinct names of the car company. So accordingly we will replace the names with the frequency of the respective car name.

```
[139]: # Calculate the frequency of each category
frequency_map = cars_df['company'].value_counts(normalize=True).to_dict()

# Map the frequencies to the 'company_name' column
cars_df['company_name_frequency_encoded'] = cars_df['company'].
    ↪map(frequency_map)

[140]: #Label Encoding
encoders_nums = {"doornumber": {"four": 4, "two": 2},
                  "cylindernumber":{"four": 4, "six": 6, "five": 5, "eight": 8,
                                   "two": 2, "twelve": 12, "three":3 }
                  }
cars_df = cars_df.replace(encoders_nums)
```

Here in this step we are creating a copy of the updated dataset i.e cars_df till this step which will be used for EDA and answering some of the questions related to this project.

```
[141]: # Created another copy of the above dataset for EDA as some additional columns
        ↪ were added prior to this step
cars_df_copy2 = cars_df.copy()
```

```
[142]: # Creating dummy variables for the below mentioned categorical variables
cars_df = pd.get_dummies(cars_df, columns=['carbody'], prefix=['body'],
        ↪ drop_first=True)
cars_df = pd.get_dummies(cars_df, columns=['fueltype'], prefix=['fuel'],
        ↪ drop_first = True)
cars_df = pd.get_dummies(cars_df, columns=['aspiration'], prefix=['asp'],
        ↪ drop_first = True)
cars_df = pd.get_dummies(cars_df, columns=['enginelocation'], prefix=['eng'],
        ↪ drop_first=True)
cars_df = pd.get_dummies(cars_df, columns=['enginetype'], prefix=['type'],
        ↪ drop_first=True)
cars_df = pd.get_dummies(cars_df, columns=['fuelsystem'], prefix=['sys'],
        ↪ drop_first=True)
cars_df = pd.get_dummies(cars_df, columns=['drivewheel'], prefix=['drive'],
        ↪ drop_first=True)
```

6.4.5 5. Feature Selections

In this stage, We will be doing the feature selection from our dataset. For this project and problem statement we will be using the **Univariate Feature Selection** where in we will be using the Correlation approach (**Pearson & Point Biserial**) for selecting the best feature with respect to with respect to the dependent variable (**Price**).

Here we will be using the two different correlation approaches (**Pearson & Point Biserial**) for feature selection. The reason for choosing two different correlation approaches is that we have mixed data in our dataset i.e numeric (**continuous & discrete**) and categorically encoded data (**binary**). Therefore relevant correlation approaches were chosen considering the type of data involved.

Once we have have calculated the correlation value with respect to price, we will setup a threshold value range say (**-0.3 to 0.3**) and will include only those features which will lie outside this range. The basic assumption behind this is that when we measure correlation of any variable with respect to dependent variable as the value is more towards the (+/-) 1, there appears to be a **strong linear relationship** with dependent variable and hence should include those variables.

For Numeric Data (**Continuous & Discrete**)

```
[143]: # Extract only numeric columns (int and float)
numeric_columns = cars_df.select_dtypes(include=['int', 'float']).columns

# Calculate Pearson correlation for each numeric column with 'price'
```

```

correlation_pc = cars_df[numeric_columns].corrwith(cars_df['price'])

# Create a DataFrame to store variable names and their correlation values
correlation_df_pc = pd.DataFrame({'Variable': correlation_pc.index, 'Pearson_
↳Correlation': correlation_pc.values})

# Filter variables based on correlation values outside the range (-0.3, 0.3)
selected_features_p_correlation = correlation_df_pc[(correlation_df_pc['Pearson_
↳Correlation'] < -0.3) | (correlation_df_pc['Pearson Correlation'] > 0.3)]

#correlation_df
selected_features_p_correlation

```

```

[143]:

```

	Variable	Pearson Correlation
2	wheelbase	0.595909
3	carlength	0.712455
4	carwidth	0.783230
6	curbweight	0.864597
7	cylindernumber	0.677018
8	enginesize	0.860063
9	boreratio	0.572685
12	horsepower	0.821715
14	citympg	-0.718290
15	highwaympg	-0.733692
16	price	1.000000
17	mileage	-0.730484
18	car_area	0.762857
19	car_volume	0.650803
20	company_name_frequency_encoded	-0.341673

For Categorical Encoded Data (**Binary**)

```

[144]: # Extract only uint8 columns
categorical_features = cars_df.select_dtypes(include='uint8').columns

# Calculate point-biserial correlation for each uint8 column with 'price'
correlation_values_pb = {}

for feature in categorical_features:
    correlation_pb, _ = pointbiserialr(cars_df[feature], cars_df['price'])
    if abs(correlation_pb) > 0.3: # Check if correlation is outside the range_
↳(-0.3, 0.3)
        correlation_values_pb[feature] = correlation_pb

selected_features_pb_correlation = pd.DataFrame(list(correlation_values_pb.
↳items()), columns=['Variable', 'Point Biserial Correlation'])

```

```
# Display the resulting correlation values
selected_features_pb_correlation
```

```
[144]:
```

	Variable	Point Biserial Correlation
0	eng_rear	0.304551
1	type_ohc	-0.338524
2	type_ohcv	0.345786
3	sys_2bbl	-0.550535
4	sys_mpfi	0.545737
5	drive_fwd	-0.636983
6	drive_rwd	0.673377

```
[145]: # Merging both the selected features into a single variable
selected_features = pd.concat([selected_features_p_correlation['Variable'],
↪selected_features_pb_correlation['Variable']], ignore_index=True).unique()
selected_features
```

```
[145]: array(['wheelbase', 'carlength', 'carwidth', 'curbweight',
'cylindernumber', 'enginesize', 'boreratio', 'horsepower',
'citympg', 'highwaympg', 'price', 'mileage', 'car_area',
'car_volume', 'company_name_frequency_encoded', 'eng_rear',
'type_ohc', 'type_ohcv', 'sys_2bbl', 'sys_mpfi', 'drive_fwd',
'drive_rwd'], dtype=object)
```

After doing the Feature Selection, here is the list of Variables which are selected for implementing different machine learning models.

Variable Names

1. wheelbase
2. carlength
3. carwidth
4. curbweight
5. cylindernumber
6. enginesize
7. boreratio
8. horsepower
9. citympg
10. highwaympg
11. mileage
12. car_area
13. eng_rear

14. type_ohc
15. type_ohcv
16. sys_2bbl
17. sys_mpf
18. drive_fwd
19. drive_rwd
20. car_volume
21. company_name_frequency_encoded

6.4.6 6. Data Splitting

Once we have selected the best features according to the respective method applied, we need to now create a separate dataset to store only those selected features with their values which will be then splitted into (**Training & Testing**) data.

So for the same, first we will merge the two dataset containing the selected features and then store it into another variable called **selected_features**.

Post this we will create a copy of original dataset with only those selected features.

```
[146]: # Merging both the selected features into a single variable
selected_features = pd.concat([selected_features_p_correlation['Variable'],
    ↪selected_features_pb_correlation['Variable']], ignore_index=True).unique()

# Create a copy of the original dataset with only the selected features
selected_features_df = cars_df[selected_features].copy()

selected_features_df
```

```
[146]:
```

	wheelbase	carlength	carwidth	curbweight	cylindernumber	enginesize	\
0	88.6	168.8	64.1	2548	4	130	
1	88.6	168.8	64.1	2548	4	130	
2	94.5	171.2	65.5	2823	6	152	
3	99.8	176.6	66.2	2337	4	109	
4	99.4	176.6	66.4	2824	5	136	
..	
200	109.1	188.8	68.9	2952	4	141	
201	109.1	188.8	68.8	3049	4	141	
202	109.1	188.8	68.9	3012	6	173	
203	109.1	188.8	68.9	3217	6	145	
204	109.1	188.8	68.9	3062	4	141	

	boreratio	horsepower	citympg	highwaympg	...	car_area	car_volume	\
0	3.47	111	21.0	27.0	...	10820.08	528019.904	
1	3.47	111	21.0	27.0	...	10820.08	528019.904	

2	2.68	154	19.0	26.0	...	11213.60	587592.640
3	3.19	102	24.0	30.0	...	11690.92	634816.956
4	3.19	115	18.0	22.0	...	11726.24	636734.832
..
200	3.78	114	23.0	28.0	...	13008.32	721961.760
201	3.78	160	19.0	25.0	...	12989.44	720913.920
202	3.58	134	18.0	23.0	...	13008.32	721961.760
203	3.01	106	26.0	27.0	...	13008.32	721961.760
204	3.78	114	19.0	25.0	...	13008.32	721961.760

	company_name_frequency_encoded	eng_rear	type_ohc	type_ohcv	sys_2bbl	\
0	0.014634	0	0	0	0	
1	0.014634	0	0	0	0	
2	0.014634	0	0	1	0	
3	0.034146	0	1	0	0	
4	0.034146	0	1	0	0	
..	
200	0.053659	0	1	0	0	
201	0.053659	0	1	0	0	
202	0.053659	0	0	1	0	
203	0.053659	0	1	0	0	
204	0.053659	0	1	0	0	

	sys_mpf	drive_fwd	drive_rwd
0	1	0	1
1	1	0	1
2	1	0	1
3	1	1	0
4	1	0	0
..
200	1	0	1
201	1	0	1
202	1	0	1
203	0	0	1
204	1	0	1

[205 rows x 22 columns]

```
[147]: # Split your data to train and test. Choose Splitting ratio wisely.

# Create a dependent Variable
dependent_variable = 'price'

# Create an Independent Variable
independent_variable = list(set(selected_features_df.columns.tolist()) -
                             {dependent_variable})
```

```

# Create variable x for storing independent variables values in it
x = selected_features_df[independent_variable].values

# Create variable y for storing dependent/target variable values in it
y = selected_features_df[dependent_variable].values

# Splitting the dataset into the Training set and Test set
x_train, x_test, y_train, y_test = train_test_split(x, y, test_size = 0.2,
↳ random_state = 0)

```

```

[148]: # Display the structure of the Training and Testing data
print(x_train.shape)
print(y_train.shape)
print(x_test.shape)
print(y_test.shape)

```

(164, 21)

(164,)

(41, 21)

(41,)

###7. Data Transformation/Scaling

Do you think that your data needs to be transformed? If yes, which transformation have you used. Explain Why? Logarithmic transformation is often employed for skewed datasets due to several advantages it offers in statistical analyses. One primary benefit is the stabilization of variance, especially when dealing with heteroscedasticity, where the variability of the data fluctuates across different levels of independent variables. Additionally, skewed distributions, whether left-skewed or right-skewed, may violate the assumption of normality in statistical methods. Logarithmic transformation helps mitigate skewness, making the distribution more symmetric and aligning it closer to a normal distribution. Moreover, when relationships between variables are not linear, as assumed in linear regression, logarithmic transformation can be used to linearize these relationships and improve model fit.

```

[149]: # Logarithmic transformation on training data
x_train_log = np.log1p(x_train)

# Logarithmic transformation on test data
x_test_log = np.log1p(x_test)

```

6.5 5. EDA Process

In this step after analysing the the dataset i.e cars_df, i have come up with certain questions related to project which will help us better understand the this project and also get a good grasp of how the price and other related variables behave with respect to each other and therefore will include bivariate anlaysis.

Q1. What are the top 10 cars by price factor?

```
[150]: # Extracting the relevant columns
columns_of_interest1 = ['CarName', 'price']
cars_data1 = cars_df_copy2[columns_of_interest1]

# Sorting the DataFrame by price in descending order
top_10_cars = cars_data1.sort_values(by='price', ascending=False).head(10)
top_10_cars
```

```
[150]:
```

	CarName	price
17	bmw x3	29575.5
74	buick regal sport coupe (turbo)	29575.5
48	jaguar xf	29575.5
129	porsche cayenne	29575.5
16	bmw x5	29575.5
15	bmw x4	29575.5
47	jaguar xj	29575.5
126	porcshce panamera	29575.5
73	buick century special	29575.5
72	buick skylark	29575.5

Q2. What are the lowest 10 cars by price factor?

```
[151]: # Extracting the relevant columns
columns_of_interest2 = ['CarName', 'price']
cars_data2 = cars_df_copy2[columns_of_interest2]

# Sorting the DataFrame by price in ascending order
lowest_10_cars = cars_data2.sort_values(by='price', ascending=True).head(10)

# Displaying the lowest 10 priced cars
lowest_10_cars
```

```
[151]:
```

	CarName	price
138	subaru	5118.0
18	chevrolet impala	5151.0
50	maxda rx3	5195.0
150	toyota corona mark ii	5348.0
76	mitsubishi mirage	5389.0
32	honda civic	5399.0
89	Nissan versa	5499.0
118	plymouth fury iii	5572.0
21	dodge rampage	5572.0
51	maxda glc deluxe	6095.0

Q3. What are the top 10 cars by car volume factor ?


```
[152]: # Extracting the relevant columns
columns_of_interest3 = ['CarName', 'car_volume']
cars_data3 = cars_df_copy2[columns_of_interest3]

# Sorting the DataFrame by price in ascending order
top_10_cars_in_volume_terms= cars_data3.sort_values(by='car_volume',
↪ascending=False).head(10)

# Displaying the highest 10 cars in terms of car volume
top_10_cars_in_volume_terms
```

```
[152]:
```

	CarName	car_volume
73	buick century special	838928.097
71	buick opel isuzu deluxe	813874.590
70	buick skyhawk	810993.618
114	peugeot 505s turbo diesel	798599.412
110	peugeot 504	798599.412
109	peugeot 504 (sw)	798599.412
68	buick century luxus (sw)	787769.849
17	bmw x3	786358.990
74	buick regal sport coupe (turbo)	784636.848
113	peugeot 504	771389.892

Q4. What are the lowest 10 cars by car volume factor ?

```
[153]: # Extracting the relevant columns
columns_of_interest4 = ['CarName', 'car_volume']
cars_data4 = cars_df_copy2[columns_of_interest4]

# Sorting the DataFrame by car_volume in ascending order
lowest_10_cars_in_volume_terms = cars_data4.sort_values(by='car_volume',
↪ascending=True).head(10)

# Displaying the lowest 10 cars in terms of car volume
lowest_10_cars_in_volume_terms
```

```
[153]:
```

	CarName	car_volume
18	chevrolet impala	452643.156
30	honda civic	469388.952
31	honda civic cvcc	469388.952
32	honda civic	504960.000
34	honda civic cvcc	504960.000
33	honda accord cvcc	504960.000
120	plymouth fury iii	507808.444
24	dodge monaco (sw)	507808.444
25	dodge colt hardtop	507808.444
26	dodge colt (sw)	507808.444

Q5. What are the top 10 cars by car area factor ?

```
[154]: # Extracting the relevant columns
columns_of_interest5 = ['CarName', 'car_area']
cars_data5 = cars_df_copy2[columns_of_interest5]

# Sorting the DataFrame by car_area in descending order
top_10_cars_in_area_terms = cars_data5.sort_values(by='car_area',
↪ascending=False).head(10)

# Displaying the highest 10 cars in terms of car area
top_10_cars_in_area_terms
```

```
[154]:
```

	CarName	car_area
73	buick century special	14795.91
70	buick skyhawk	14404.86
71	buick opel isuzu deluxe	14404.86
74	buick regal sport coupe (turbo)	14163.12
17	bmw x3	13967.30
48	jaguar xf	13892.16
47	jaguar xj	13892.16
6	audi 100ls	13700.97
7	audi 5000	13700.97
8	audi 4000	13700.97

Q6. What are the lowest 10 car by car area factor ?

```
[155]: # Extracting the relevant columns
columns_of_interest6 = ['CarName', 'car_area']
cars_data6 = cars_df_copy2[columns_of_interest6]

# Sorting the DataFrame by car_area in ascending order
lowest_10_cars_in_area_terms = cars_data6.sort_values(by='car_area',
↪ascending=True).head(10)

# Displaying the lowest 10 cars in terms of car area
lowest_10_cars_in_area_terms
```

```
[155]:
```

	CarName	car_area
18	chevrolet impala	8508.33
31	honda civic cvcc	9239.94
30	honda civic	9239.94
34	honda civic cvcc	9600.00
33	honda accord cvcc	9600.00
32	honda civic	9600.00
45	isuzu D-Max V-Cross	9915.24
44	isuzu D-Max	9915.24
19	chevrolet monte carlo	9915.24

138 subaru 9947.46

Q7. What is the average car volume against each car company ?

```
[156]: # Extracting relevant columns
columns_of_interest7 = ['company', 'car_volume']

# Creating a DataFrame with the relevant columns
average_car_volume_by_company = cars_df_copy2[columns_of_interest7]

# Grouping by 'company' and calculating the average volume
average_car_volume_by_company = average_car_volume_by_company.
    ↳groupby('company')['car_volume'].mean().reset_index().sort_values(by =_
    ↳'car_volume', ascending = True)

# Renaming the columns for clarity
average_car_volume_by_company.columns = ['company', 'average car volume']

# Displaying the DataFrame
average_car_volume_by_company
```

```
[156]:
```

	company	average car volume
4	Chevrolet	497806.332000
5	Dodge	534625.439111
7	Isuzu	543500.324500
0	Alfa-Romero	547877.482667
6	Honda	551744.387462
14	Plymouth	551834.287429
11	Mitsubishi	556000.344615
15	Porsche	581886.421800
18	Subaru	589594.428333
12	Nissan	597112.215778
9	Mazda	598388.211647
19	Toyota	601358.206250
20	Volkswagen	625175.615500
16	Renault	630440.820000
10	Mercury	664789.760000
2	Bmw	673741.664000
1	Audi	688157.514000
17	Saab	696139.290000
8	Jaguar	704646.084000
21	Volvo	721492.677818
13	Peugeot	747793.413091
3	Buick	770478.162750

Q8. What is the average car area against each car company ?

```
[157]: # Extracting relevant columns for car_area
columns_of_interest8 = ['company', 'car_area']

# Creating a DataFrame with the relevant columns for car_area
average_car_area_by_company = cars_df_copy2[columns_of_interest8]

# Grouping by 'company' and calculating the average car area
average_car_area_by_company = average_car_area_by_company.
↳groupby('company')['car_area'].mean().reset_index().sort_values(by_
↳='car_area', ascending = True)

# Renaming the columns for clarity
average_car_area_by_company.columns = ['company', 'average car area']

# Displaying the DataFrame for car_area
average_car_area_by_company
```

```
[157]:
```

	company	average car area
4	Chevrolet	9507.750000
5	Dodge	10334.722222
6	Honda	10354.879231
7	Isuzu	10408.315000
14	Plymouth	10602.431429
0	Alfa-Romero	10951.253333
11	Mitsubishi	10969.524615
18	Subaru	10972.526667
12	Nissan	11137.121111
19	Toyota	11199.268125
9	Mazda	11208.834118
20	Volkswagen	11326.371667
15	Porsche	11392.728000
16	Renault	11922.315000
10	Mercury	12131.200000
2	Bmw	12280.135000
17	Saab	12408.900000
1	Audi	12624.977143
21	Volvo	12831.534545
13	Peugeot	13072.030000
8	Jaguar	13772.780000
3	Buick	13812.711250

Q9. What is the average mileage across each car company ?

```
[158]: # Extracting relevant columns for mileage and company
columns_of_interest9 = ['company', 'mileage']
average_mileage_by_company = cars_df_copy2[columns_of_interest9]
```

```

# Grouping by 'company' and calculating the average mileage
average_mileage_by_company = average_mileage_by_company.
↳groupby('company')['mileage'].mean().reset_index().sort_values(by =
↳'mileage', ascending = True)

# Renaming the columns for clarity
average_mileage_by_company.columns = ['company', 'average mileage']

# Displaying the DataFrame
average_mileage_by_company

```

```

[158]:
      company  average mileage
8      Jaguar      15.933333
3      Buick      19.500000
15     Porsche      20.840000
1      Audi      20.971429
10     Mercury      21.000000
2      Bmw      21.775000
0  Alfa-Romero      22.866667
21     Volvo      23.036364
17     Saab      23.133333
13     Peugeot      24.127273
16     Renault      26.200000
11  Mitsubishi      27.415385
18     Subaru      28.100000
9      Mazda      28.200000
12     Nissan      29.322222
19     Toyota      29.662500
5      Dodge      30.444444
14     Plymouth      30.542857
20  Volkswagen      31.116667
6      Honda      32.100000
7      Isuzu      33.000000
4      Chevrolet      42.300000

```

Q10. What is the average mileage of against different fuel types i.e (gas vs diseal ?)

```

[159]: # Extracting relevant columns
columns_of_interest10 = ['fueltype', 'mileage']
average_mileage_by_fueltype = cars_df_copy2[columns_of_interest10]

# Grouping by 'fueltype' and calculating the average mileage
average_mileage_by_fueltype = average_mileage_by_fueltype.
↳groupby('fueltype')['mileage'].mean().reset_index()

# Renaming the columns for clarity
average_mileage_by_fueltype.columns = ['fueltype', 'average mileage']

```

```
# Displaying the DataFrame
average_mileage_by_fueltype
```

```
[159]: fueltype average mileage
0    diesel      32.030000
1      gas      26.894054
```

Q11. What is the average price of cars against different risk rating ?

```
[160]: # Group by 'symboling' and calculate the average price for each group
average_price_against_safety_ratings = pd.DataFrame(cars_df_copy2.
↳groupby('symboling')['price'].mean()).reset_index().sort_values(by='price',
↳ascending=True)

# Rename the 'price' column to 'average price'
average_price_against_safety_ratings = average_price_against_safety_ratings.
↳rename(columns={'price': 'average price'})

# Print or use the DataFrame
average_price_against_safety_ratings
```

```
[160]: symboling average price
3         1    9711.064815
4         2   10109.281250
2         0   13670.151746
0        -2   15781.666667
5         3   16468.037037
1        -1   17029.181818
```

Q12. What is the average price of cars for each car company ?

```
[161]: # Group by 'company' and calculate the average price for each group
average_price_against_company = pd.DataFrame(cars_df_copy2.
↳groupby('company')['price'].mean()).reset_index().sort_values(by='price',
↳ascending=True)

# Rename the 'price' column to 'average price'
average_price_against_company = average_price_against_company.
↳rename(columns={'price': 'average price'})

# Print or use the DataFrame
average_price_against_company
```

```
[161]: company average price
4    Chevrolet    6007.000000
5      Dodge    7875.444444
```

14	Plymouth	7963.428571
6	Honda	8184.692308
18	Subaru	8541.250000
7	Isuzu	8916.500000
11	Mitsubishi	9239.769231
16	Renault	9595.000000
19	Toyota	9885.812500
20	Volkswagen	10077.500000
12	Nissan	10415.666667
9	Mazda	10652.882353
17	Saab	15223.333333
13	Peugeot	15489.090909
0	Alfa-Romero	15498.333333
10	Mercury	16503.000000
1	Audi	17859.166714
21	Volvo	18063.181818
2	Bmw	23590.187500
15	Porsche	28064.000000
3	Buick	28731.687500
8	Jaguar	29575.500000

Q13. What is the average price of cars against different car body types ?

```
[162]: # Group by 'carbody' and calculate the average price for each group
average_price_against_body_type = pd.DataFrame(cars_df_copy2.
↳groupby('carbody')['price'].mean()).reset_index().sort_values(by='price',
↳ascending=True)

# Rename the 'price' column to 'average price'
average_price_against_body_type = average_price_against_body_type.
↳rename(columns={'price': 'average price'})

# Print or use the DataFrame
average_price_against_body_type
```

```
[162]:      carbody  average price
2   hatchback    10350.580957
4      wagon    12371.960000
3      sedan    13788.859375
1   hardtop    19304.812500
0 convertible    19735.000000
```

Q14. What is the average car price against different engine types of the car ?

```
[163]: # Group by 'enginetype' and calculate the average price for each group
```

```

average_price_against_engine_type = pd.DataFrame(cars_df_copy2.
↳groupby('enginetype')['price'].mean()).reset_index().sort_values(by='price',
↳ascending=True)

# Rename the 'price' column to 'average price'
average_price_against_engine_type = average_price_against_engine_type.
↳rename(columns={'price': 'average price'})

# Print or use the DataFrame
average_price_against_engine_type

```

```

[163]:  enginetype  average price
3      ohc      11423.690318
4      ohcf     12748.100000
6      rotor    13020.000000
2       l      14627.583333
0      dohc     17395.666667
5      ohcv     21735.115385
1      dohcv    29575.500000

```

Q15. What is the average car price against different fuel systems of cars ?

```

[164]: # Group by 'fuelsystem' and calculate the average price for each group
average_price_against_fuel_system = pd.DataFrame(cars_df_copy2.
↳groupby('fuelsystem')['price'].mean()).reset_index().sort_values(by='price',
↳ascending=True)

# Rename the 'price' column to 'average price'
average_price_against_fuel_system = average_price_against_fuel_system.
↳rename(columns={'price': 'average price'})

# Print or use the DataFrame
average_price_against_fuel_system

```

```

[164]:  fuelsystem  average price
1      2bb1      7478.151515
0      1bb1      7555.545455
6      spdi     10990.444444
7      spfi     11048.000000
2      4bb1     12145.000000
4       mfi     12964.000000
3       idi     15736.925000
5      mpfi     16804.789011

```

Q16. What is the average price of cars against fuel type i.e (gas vs diesel) ?


```
[165]: # Group by 'fueltype' and calculate the average price for each group
average_price_against_fuel_type = pd.DataFrame(cars_df_copy2.
    ↳groupby('fueltype')['price'].mean()).reset_index().sort_values(by='price',
    ↳ascending=True)

# Rename the 'price' column to 'average price'
average_price_against_fuel_type = average_price_against_fuel_type.
    ↳rename(columns={'price': 'average price'})

# Print or use the DataFrame
average_price_against_fuel_type
```

```
[165]: fueltype average price
1      gas  12517.190092
0     diesel  15736.925000
```


6.5.1 Visualizing The Above Data.

```
[166]: # Plotting
plt.figure(figsize=(10, 6), dpi=110) # Setting the figure size & dpi
sns.set_style('darkgrid') # Setting the plot style to darkgrid using Seaborn

# Creating a bar plot using Seaborn
# x-axis: 'symboling', y-axis: 'price', data source: 'cars_df_copy2'
# Additional parameters:
#   - edgecolor: Black border for bars
#   - hue: 'fueltype', differentiating bars by fuel type
#   - errorbar: None, no error bars displayed
#   - palette: 'Set1', color palette for the plot
sns.barplot(x='symboling', y='price', data=cars_df_copy2, edgecolor='black',
    ↳hue='fueltype', errorbar=None, palette='Set1')

# Setting plot title, x-axis label, and y-axis label
plt.title('Average Price of Cars Against Safety Ratings')
plt.xlabel('Safety Rating')
plt.ylabel('Average Price of Cars')

# Displaying the plot
plt.show()
```



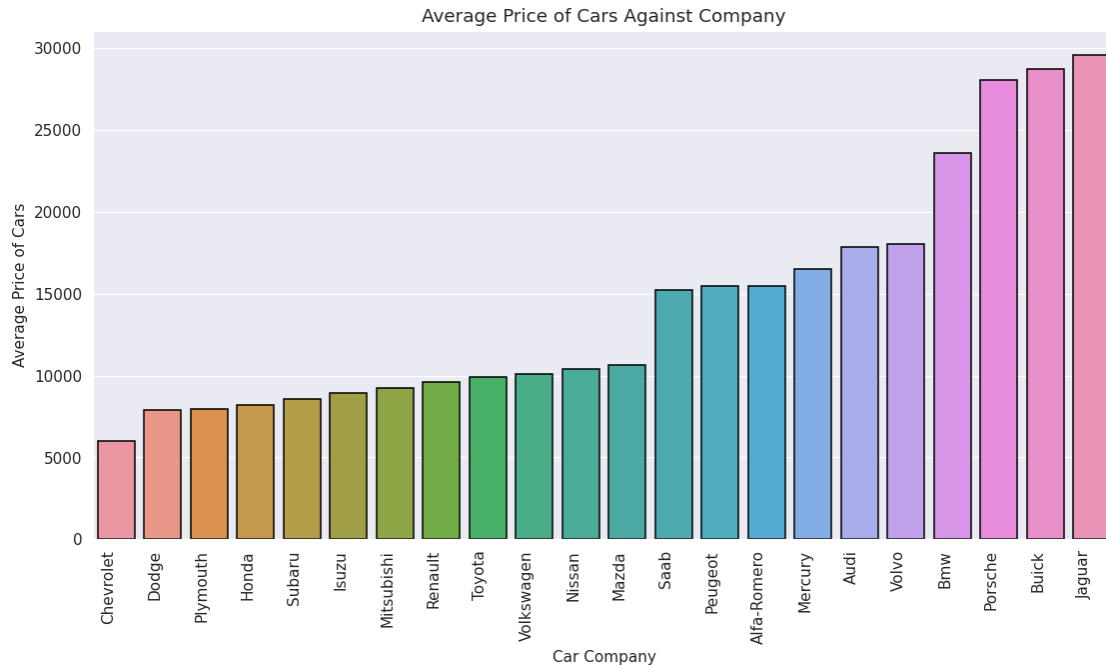
```
[167]: # Plotting
plt.figure(figsize=(12, 6), dpi=110) # Setting the figure size & dpi
sns.set_style('darkgrid') # Setting the plot style to darkgrid using Seaborn

# Creating a bar plot using Seaborn
# x-axis: 'company', y-axis: 'price', data source:
↳ 'average_price_against_company'
# Additional parameters:
# - edgecolor: Black border for bars
sns.barplot(x='company', y='average price', data=average_price_against_company,
↳ edgecolor='black')

# Setting plot title, x-axis label, and y-axis label
plt.title('Average Price of Cars Against Company')
plt.xlabel('Car Company')
plt.ylabel('Average Price of Cars')

# Rotating x-axis labels for better visibility
plt.xticks(rotation=90, ha='right')


# Displaying the plot
plt.show()
```



```
[168]: # Plotting
plt.figure(figsize=(12, 6)) # Setting the figure size
# Creating a bar plot using Seaborn
# x-axis: 'average mileage', y-axis: 'company', data source:
↳ 'average_mileage_by_company'
# Additional parameter:
# - palette: 'viridis', color palette for the plot
sns.barplot(x='average mileage', y='company', data=average_mileage_by_company,
↳ palette='viridis')

# Setting plot title, x-axis label, and y-axis label
plt.title('Average Mileage by Company')
plt.xlabel('Average Mileage')
plt.ylabel('Company')

# Displaying the plot
plt.show()
```




```
[169]: # Plotting
plt.figure(figsize=(10, 6), dpi=110) # Setting the figure size
sns.set_style('darkgrid') # Setting the plot style to darkgrid using Seaborn

# Creating a bar plot using Seaborn
# x-axis: 'carbody', y-axis: 'price', data source: 'cars_df_copy1'
# Additional parameters:
# - hue: 'fueltype', differentiating bars by fuel type
# - errorbar: None, no error bars displayed
# - edgecolor: Black border for bars
sns.barplot(x='carbody', y='price', data=cars_df_copy2, hue='fueltype',
            errorbar=None, edgecolor='black', palette='bright')

# Setting plot title, x-axis label, and y-axis label
plt.title('Average Price of Cars Against Car Body Type')
plt.xlabel('Car Body Type')
plt.ylabel('Average Price of Cars')

# Displaying the plot
plt.show()
```




```
[170]: # Plotting
plt.figure(figsize=(12, 6)) # Setting the figure size
# Creating a bar plot using Seaborn
# x-axis: 'enginetype', y-axis: 'average price', data source:
↳ 'average_price_against_engine_type'
# Additional parameters:
# - edgecolor: Black border for bars
# - palette: 'bright', bright color palette for the plot
sns.barplot(x='enginetype', y='average price',
↳ data=average_price_against_engine_type, edgecolor='black', palette='bright')

# Setting plot title, x-axis label, and y-axis label
plt.title('Average Price Against Engine Type')
plt.xlabel('Engine Type')
plt.ylabel('Average Price')

# Rotating x-axis labels for better visibility
plt.xticks(rotation=45, ha='right')

# Displaying the plot
plt.show()
```




```
[171]: # Plotting
plt.figure(figsize=(12, 6)) # Setting the figure size
# Creating a bar plot using Seaborn
# x-axis: 'fuelsystem', y-axis: 'average price', data source:
# ↪ 'average_price_against_fuel_system'
# Additional parameters:
# - edgecolor: Black border for bars
# - palette: 'bright', bright color palette for the plot
sns.barplot(x='fuelsystem', y='average price',
# ↪ data=average_price_against_fuel_system, edgecolor='black', palette='bright')

# Setting plot title, x-axis label, and y-axis label
plt.title('Average Price Against Fuel System')
plt.xlabel('Fuel System')
plt.ylabel('Average Price')

# Rotating x-axis labels for better visibility
plt.xticks(rotation=45, ha='right')


# Displaying the plot
plt.show()
```



```
[172]: # Plotting
plt.figure(figsize=(8, 5)) # Setting the figure size
# Creating a bar plot using Seaborn
# x-axis: 'fueltype', y-axis: 'average price', data source:
#   ↳ 'average_price_against_fuel_type'
# Additional parameters:
#   - edgecolor: Black border for bars
#   - palette: 'bright', bright color palette for the plot
sns.barplot(x='fueltype', y='average price',
#   ↳ data=average_price_against_fuel_type, edgecolor='black', palette='bright')

# Setting plot title, x-axis label, and y-axis label
plt.title('Average Price Against Fuel Type')
plt.xlabel('Fuel Type')
plt.ylabel('Average Price')

# Displaying the plot
plt.show()
```



6.6 6. ML Model Implementation

6.6.1 ML Model 1 - Linear Regression

```
[173]: # Create a object to run Linear Regression
regressor = LinearRegression()

# Fit into the Algorithm
regressor.fit(x_train,y_train)
```

```
[173]: LinearRegression()
```

```
[174]: # Calculate and show the Coefficients of Independent variables
regressor.coef_
```

```
[174]: array([-2.96366573e+03,  1.04938440e+03, -2.89222273e+01, -1.41493439e+03,
        2.01363127e+01,  8.88570310e+02, -1.08973681e+02,  5.25129440e+02,
        3.52717515e+00, -1.08474287e+04, -8.15859597e+02, -4.66495226e+02,
        2.15892426e+01,  9.11549531e+01, -2.63051185e+03, -1.21164820e+02,
        2.46858718e+01,  2.93855733e-03, -1.47523782e+03,  9.55096440e+03,
        1.79489569e+03])
```

```
[175]: # Calculate and display the Intercept Value
regressor.intercept_
```


[175]: 203808.83025884215

```
[176]: # Calculate and display the Predicted values for Training dataset
y_predict_train = regressor.predict(x_train)

# Calculate and display the Predicted value for the Test dataset
y_predict_test = regressor.predict(x_test)
```

```
[177]: # Display the Predicted values for Training Data
y_predict_train
```

```
[177]: array([16183.6130289 , 28069.53838862, 15209.20296173, 10646.00170657,
        7880.72751545,  7139.19172011, 14625.43118881, 10832.65106574,
        10289.83059088, 22254.83151042,  6044.84521341,  8802.73580508,
        10107.5576675 , 19316.7679723 , 18785.15704832,  9971.03356385,
        18695.26299279,  8790.57676719,  7061.59386682,  6395.00436652,
        13913.85107723,  7604.73159651, 10258.27736287,  9354.5708033 ,
        14733.3127523 ,  9856.18023438,  6414.53137363, 10064.64102861,
        6313.78114488,  6621.93675189,  8722.47882443,  6663.32029003,
        6034.67829072,  6509.48540278, 17214.06817249,  6246.77961549,
        16945.94337603,  7179.84570626, 17793.4458281 , 19638.50153521,
        20740.62200205,  9884.51372636, 17020.07353927, 10067.61855647,
        29523.76809782, 14179.22235905, 14950.39695877, 18491.46150522,
        6393.84108988, 10258.27736287, 17209.7273977 , 10107.5576675 ,
        9372.37302937, 17190.37315623,  9775.44985511, 12396.20026444,
        17895.89808234, 29543.69590315, 16969.46450804,  8711.89729899,
        5029.70606961,  6347.28695267, 11487.98399395,  6267.92786794,
        6481.28505975,  6009.23389441, 19138.77390472,  6197.27312031,
        13934.65931811,  6176.2361125 , 17214.06817249,  6402.37566697,
        10311.18499011,  8311.1828302 , 11038.60763469, 24049.01996631,
        32512.08631998, 28252.11088357,  9715.09322841, 18521.76288136,
        17716.89392304,  6213.84829943,  9668.19981142,  9365.87575679,
        14278.68281474,  6444.77604354, 23456.91137593,  7309.19426991,
        9135.80709983,  9960.05000222, 21293.1930906 , 20969.30489244,
        25086.50759671, 14643.06706455,  6970.75116071,  6270.60281237,
        8920.31021215,  6640.31226545,  6059.83933258, 10096.78616181,
        12396.20026444,  9248.8989446 , 15569.74932581,  7302.17804315,
        9372.20667905, 14685.3352563 , 13934.65931811,  7167.42243952,
        8926.18693531,  9982.89070123, 19864.29335055, 13561.1335623 ,
        10217.66715897, 29240.42041201,  9559.83557852, 13691.63904282,
        6678.062767 ,  7004.15665549, 29678.96380439,  7705.21328092,
        9281.29683484, 17987.44046131,  5649.42270966,  8333.33973766,
        13599.21339291, 29523.76809782, 18869.45645149,  6268.45446072,
        13613.32209351,  6855.17210469, 19146.00065405,  9294.48501326,
        14434.34668961,  9213.40495311, 14329.79215384, 10062.4714524 ,
        9076.19057822, 17139.93800924,  8923.33942081, 17395.33954422,
        26285.30992536,  7097.56376095,  7139.19172011,  9836.598872 ,
```

```

9515.59460462, 19912.10029624, 8885.2444019 , 17014.66649686,
9662.50756476, 11360.78162836, 7491.7683252 , 10942.23212102,
27041.2324798 , 10942.23212102, 7548.31789598, 6227.31513078,
19435.969594 , 14687.36670671, 17493.44916029, 24170.8617244 ,
11050.16414874, 18342.54853549, 28252.11088357, 15181.60924488])

```

```

[178]: # Display the Predicted values for Test Data
y_predict_test

```

```

[178]: array([ 6680.95616577, 18819.46825352, 14102.66201594, 2318.71376799,
10647.2095531 , 14103.59943926, 6486.65329014, 7172.35740617,
19027.55571255, 7648.77847853, 17522.89928983, 28159.42284203,
9354.5708033 , 13077.2998493 , 6485.06821751, 13652.12102015,
13636.92277518, 21357.29415886, 9372.89529561, 5561.24333092,
10756.84219783, 17099.66538382, 12994.53817043, 14227.05056949,
19905.91555845, 5144.11478822, 6178.13613344, 17969.66335559,
6330.3519372 , 5957.55125325, 10057.03703102, 12522.41867227,
18711.22880055, 9692.91283857, 6130.38283556, 25354.31349701,
9859.1661766 , 14141.12298392, 6746.51904771, 27935.71059694,
6361.61708728])

```

```

[179]: # Create a DataFrame to display the Actual vs Predicted values for the Test
↳ dataset for Linear Regression
df_actual_vs_predicted_linear_model = pd.DataFrame({'Actual Values': y_test,
↳ 'Predicted Values': y_predict_test})
df_actual_vs_predicted_linear_model

```

```

[179]:
Actual Values Predicted Values
0          6795.0      6680.956166
1         15750.0     18819.468254
2         15250.0     14102.662016
3          5151.0      2318.713768
4          9995.0     10647.209553
5         11199.0     14103.599439
6          5389.0      6486.653290
7          7898.0      7172.357406
8         17199.0     19027.555713
9          6529.0      7648.778479
10        20970.0     17522.899290
11        29575.5     28159.422842
12        10945.0      9354.570803
13        18344.0     13077.299849
14         8916.5      6485.068218
15         9989.0     13652.121020
16         9295.0     13636.922775
17        18920.0     21357.294159
18         7895.0      9372.895296

```

19	6488.0	5561.243331
20	9959.0	10756.842198
21	15580.0	17099.665384
22	9895.0	12994.538170
23	11549.0	14227.050569
24	15998.0	19905.915558
25	5118.0	5144.114788
26	6938.0	6178.136133
27	16695.0	17969.663356
28	8358.0	6330.351937
29	5499.0	5957.551253
30	7975.0	10057.037031
31	12290.0	12522.418672
32	22018.0	18711.228801
33	8948.0	9692.912839
34	6849.0	6130.382836
35	29575.5	25354.313497
36	11595.0	9859.166177
37	18150.0	14141.122984
38	6377.0	6746.519048
39	29575.5	27935.710597
40	8916.5	6361.617087

```
[180]: # Calculate the  $r^2$  score for Training data & display
r_squared_train_for_linear_regression = round(r2_score(y_train,
    ↪ y_predict_train),4)
print('R Squared for Training data :', r_squared_train_for_linear_regression)

# Calculate the  $r^2$  score for Test data & display
r_squared_test_for_linear_regression = round(r2_score(y_test, y_predict_test),4)
print('R Squared for Test data :', r_squared_test_for_linear_regression)

# Number of observations (n)
n = len(y_test)

# Number of predictors (k) - you need to replace this with the actual number of
    ↪ predictors in your model
k = 19

# Calculate adjusted R Squared for Test data & display
adjusted_r_squared_test_for_linear_regression = round(1 - ((1 -
    ↪ r_squared_test_for_linear_regression) * (n - 1) / (n - k - 1)),4)

print("Adjusted R-squared for test data:",
    ↪ adjusted_r_squared_test_for_linear_regression)
```

R Squared for Training data : 0.9103

R Squared for Test data : 0.8723
Adjusted R-squared for test data: 0.7568

```
[181]: # Calculating Mean Squared Error (MSE) for Linear Regression
mse_lr = round(mean_squared_error(y_test, y_predict_test), 4)

# Calculating Root Mean Squared Error (RSME) for Linear Regression
rsme_lr = round(math.sqrt(mean_squared_error(y_test, y_predict_test)), 4)


# Calculating Mean Absolute Error (MAE) for Linear Regression
mae_lr = round(mean_absolute_error(y_test, y_predict_test), 4)

# Calculating Mean Absolute Percentage Error (MAPE) for Linear Regression
mape_lr = round(mean_absolute_percentage_error(y_test, y_predict_test), 4)
```

```
[182]: # Performance of Linear Regression Model
print('Performance of Linear Regression Model')
print("MSE :", round(mean_squared_error(y_test, y_predict_test), 4))
print("RMSE :", round(math.sqrt(mean_squared_error(y_test, y_predict_test)), 4))
print('MAE:', round(mean_absolute_error(y_test, y_predict_test), 4))
print('MAPE:', round(mean_absolute_percentage_error(y_test, y_predict_test), 4))
```

Performance of Linear Regression Model
MSE : 5596513.4284
RMSE : 2365.6951
MAE: 1967.1952
MAPE: 0.1698

```
[183]: # Plot the lineplot for Linear Regression
plt.figure(figsize=(12, 6), dpi = 100)
sns.set_style('darkgrid')
sns.lineplot(data=df_actual_vs_predicted_linear_model, palette="tab10",
             linewidth=2.5)
plt.title('Actual vs Predicted values for Linear Regression')
plt.xlabel('Observation')
plt.ylabel('Values')
plt.show()
```



6.6.2 ML Model - 2 Lasso Regression

Cross- Validation & Hyperparameter Tuning (Lasso Regression)

```
[184]: # Creating an object to run Lasso Regression
lasso = Lasso(max_iter = 6000)

min_alpha = 0.001
max_alpha = 0.2
num_vals = 10

parameters = {'alpha': np.logspace(np.log10(min_alpha), np.log10(max_alpha), num_vals)}

# Applying the GridSearchCV with cross validation and hyperparameter tuning
lasso_cv_hype = GridSearchCV(lasso, parameters, cv = 5, scoring = "neg_mean_squared_error")

# Fit into the Algorithm
lasso_cv_hype.fit(x_train, y_train)
```

```
/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 2.980e+08, tolerance: 6.361e+05
  model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
```

```

packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 2.034e+08, tolerance: 6.178e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 2.261e+08, tolerance: 5.247e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 2.484e+08, tolerance: 5.872e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 2.634e+08, tolerance: 5.900e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 2.980e+08, tolerance: 6.361e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 2.034e+08, tolerance: 6.178e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 2.261e+08, tolerance: 5.247e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 2.484e+08, tolerance: 5.872e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-

```

```

packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 2.634e+08, tolerance: 5.900e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 2.980e+08, tolerance: 6.361e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 2.034e+08, tolerance: 6.178e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 2.261e+08, tolerance: 5.247e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 2.484e+08, tolerance: 5.872e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 2.634e+08, tolerance: 5.900e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 2.980e+08, tolerance: 6.361e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 2.034e+08, tolerance: 6.178e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-

```

```

packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 2.261e+08, tolerance: 5.247e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 2.484e+08, tolerance: 5.872e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 2.635e+08, tolerance: 5.900e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 2.981e+08, tolerance: 6.361e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 2.034e+08, tolerance: 6.178e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 2.261e+08, tolerance: 5.247e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 2.484e+08, tolerance: 5.872e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 2.635e+08, tolerance: 5.900e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-

```



```

packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 2.981e+08, tolerance: 6.361e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 2.035e+08, tolerance: 6.178e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 2.262e+08, tolerance: 5.247e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 2.485e+08, tolerance: 5.872e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 2.635e+08, tolerance: 5.900e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 2.981e+08, tolerance: 6.361e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 2.036e+08, tolerance: 6.178e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 2.262e+08, tolerance: 5.247e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-

```

```

packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 2.485e+08, tolerance: 5.872e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 2.636e+08, tolerance: 5.900e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 2.982e+08, tolerance: 6.361e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 2.037e+08, tolerance: 6.178e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 2.263e+08, tolerance: 5.247e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 2.486e+08, tolerance: 5.872e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 2.637e+08, tolerance: 5.900e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 2.984e+08, tolerance: 6.361e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-

```

```

packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 2.039e+08, tolerance: 6.178e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 2.265e+08, tolerance: 5.247e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 2.488e+08, tolerance: 5.872e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 2.639e+08, tolerance: 5.900e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 2.986e+08, tolerance: 6.361e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 2.043e+08, tolerance: 6.178e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 2.269e+08, tolerance: 5.247e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 2.491e+08, tolerance: 5.872e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-

```

```
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 2.642e+08, tolerance: 5.900e+05
```

```
model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
```

```
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 3.333e+08, tolerance: 7.399e+05
```

```
model = cd_fast.enet_coordinate_descent(
```

```
[184]: GridSearchCV(cv=5, estimator=Lasso(max_iter=6000),
           param_grid={'alpha': array([0.001          , 0.00180165, 0.00324594,
0.00584804, 0.0105361 ,
           0.01898235, 0.03419952, 0.0616155 , 0.11100946, 0.2          ])}),
           scoring='neg_mean_squared_error')
```

```
[185]: # First fit the best alpha value into the model
lasso_alpha = Lasso(alpha = lasso_cv_hype.best_params_['alpha'])
lasso_alpha.fit(x_train,y_train)
```

```
/usr/local/lib/python3.10/dist-
```

```
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 3.421e+08, tolerance: 7.399e+05
```

```
model = cd_fast.enet_coordinate_descent(
```

```
[185]: Lasso(alpha=0.20000000000000004)
```

```
[186]: # Calculate and show the Coefficients
print('Coefficients :', lasso_alpha.coef_)
```

```
Coefficients : [ 2.35579445e+02  1.03038506e+03 -2.92453766e+02 -2.75391196e+02
 2.18739705e+01  1.09608301e+03  1.01826338e+02  5.10714375e+02
 3.94187491e+00 -1.29633380e+04 -5.42389285e+02 -2.79123974e+02
 3.83184875e+00  1.56483885e+02 -2.82519292e+03 -6.84427862e+01
 1.84083037e+01  3.33364696e-03 -1.70132558e+03  9.76927014e+03
 2.11667880e+03]
```

```
[187]: # Calculate and show the Intercept value
print('Intercept :', lasso_alpha.intercept_)
```

```
Intercept : -7634.513956755023
```

```
[188]: # Calculate Predicted values for Training dataset
y_train_pred_for_lasso_cv_hype = lasso_alpha.predict(x_train)
```

```
# Calculate the Predicted values for Test data
y_test_pred_for_lasso_cv_hype = lasso_alpha.predict(x_test)
```

```
[189]: # Show the Predicted Values for Training
print('Predicted Valu/es for Training Data :', y_train_pred_for_lasso_cv_hype)
```

```
Predicted Valu/es for Training Data : [16472.44187689 27066.45118204
15276.5119111 10571.76895374
 7860.94049905 6985.40025404 14695.33926603 10969.09674638
10419.55217016 21747.148711 5992.02416909 8675.11940083
10312.67060528 19899.73432115 19063.63059583 10112.45735868
18801.47925423 8877.43102097 6898.67900602 5907.04848234
13776.22195082 7427.01637775 10482.16635288 9265.94110837
14860.92553034 9992.87361874 6383.04726947 10055.34232855
 6200.67369161 6492.70675843 8966.54443178 6517.80102547
 5626.49867214 6812.0757814 17473.2702254 6071.41554903
17239.76187051 7493.30795629 17734.92018304 19986.24528715
21110.85734677 9559.9016493 17256.46710535 10171.21405083
29506.6849237 14429.61323854 15278.22096484 18441.95871
 6306.61580563 10482.16635288 17660.56331519 10312.67060528
 9931.09188212 17327.80900485 9473.15522059 12660.87457242
17722.63309009 27888.64512174 16969.69449794 8954.71880705
 4821.62185299 6190.80155414 11790.80675783 6149.42931778
 6515.32145449 5909.50175108 19297.95297701 6275.51154675
13440.61102348 5992.57805083 17473.2702254 6248.03897594
10541.29447653 8496.84158107 11192.02227371 24143.24534966
33125.22854393 27909.83331498 9835.19862233 18552.58721213
18210.92005865 6144.12507522 9318.2546528 9100.63980507
14258.64937382 6461.09721322 23575.22033678 7354.07366184
 9252.1995691 9808.02466169 22116.22921661 20168.04306222
25115.87153384 14715.04864058 6917.41533759 6155.63297617
 9375.32156481 6540.72061533 5862.4961788 10325.12924647
12660.87457242 8934.01186812 15669.92691518 7323.10024187
 9285.65048292 14951.11506968 13440.61102348 7025.15059662
 8813.08502268 10449.97913106 20250.34391733 13382.03445981
10459.64341348 28292.12393769 9715.86404293 13527.88383149
 6558.06366292 6673.52260958 29680.12741975 7489.30962915
 9821.63354733 17951.72330309 5293.5062814 7484.85788115
13354.68855525 29506.6849237 19262.2510301 6157.78669457
13370.45605489 7126.0343244 19051.64150091 9377.47869829
14497.66295131 9338.92081712 14365.46752975 10286.20091744
 8669.51892016 17456.56499056 8875.29629586 17446.45884617
27006.9520704 7117.56116724 6985.40025404 9670.05903983
 9598.24758621 19907.20339282 9324.6911679 16952.67729679
 9695.13060886 11251.12442879 6921.67909603 11030.21106306
25720.43355531 11030.21106306 7518.53962812 6223.29111766
```

```
19828.8809122 14669.93621242 17488.56976867 24091.10635689
11539.78823027 18598.63176246 27909.83331498 15123.70433257]
```

```
[190]: # Show the Predicted Values for Test Data
print('Predicted Values for Test Data:', y_test_pred_for_lasso_cv_hype)
```

```
Predicted Values for Test Data: [ 6537.51040002 19230.65141305 14440.60621848
-409.16184425
10840.84976471 13918.37666739 6691.70402521 6823.0049023
19367.16007484 7426.23963059 17994.1169386 29248.39625902
9265.94110837 13549.97146082 6457.77709602 13413.8166789
13406.5150757 20601.64930233 9506.9446727 5194.95940864
10809.02752658 17412.81704955 12886.67962628 14056.34228924
20217.56620533 4737.18578431 6052.40096379 17804.98133546
6235.66205725 5748.18180641 10159.3884261 12632.64280894
19156.5323485 9909.95716837 5941.333677 25061.17625389
9671.36544475 14104.91625233 6546.65951204 27158.63275271
6319.81147417]
```

```
[191]: # Create a DataFrame to display the Actual vs Predicted values for the Test
↳dataset for Lasso Regression
df_actual_vs_predicted_lasso_model = pd.DataFrame({'Actual Values': y_test,
↳'Predicted Values': y_test_pred_for_lasso_cv_hype})
df_actual_vs_predicted_lasso_model
```

```
[191]:
```

	Actual Values	Predicted Values
0	6795.0	6537.510400
1	15750.0	19230.651413
2	15250.0	14440.606218
3	5151.0	-409.161844
4	9995.0	10840.849765
5	11199.0	13918.376667
6	5389.0	6691.704025
7	7898.0	6823.004902
8	17199.0	19367.160075
9	6529.0	7426.239631
10	20970.0	17994.116939
11	29575.5	29248.396259
12	10945.0	9265.941108
13	18344.0	13549.971461
14	8916.5	6457.777096
15	9989.0	13413.816679
16	9295.0	13406.515076
17	18920.0	20601.649302
18	7895.0	9506.944673
19	6488.0	5194.959409
20	9959.0	10809.027527

21	15580.0	17412.817050
22	9895.0	12886.679626
23	11549.0	14056.342289
24	15998.0	20217.566205
25	5118.0	4737.185784
26	6938.0	6052.400964
27	16695.0	17804.981335
28	8358.0	6235.662057
29	5499.0	5748.181806
30	7975.0	10159.388426
31	12290.0	12632.642809
32	22018.0	19156.532348
33	8948.0	9909.957168
34	6849.0	5941.333677
35	29575.5	25061.176254
36	11595.0	9671.365445
37	18150.0	14104.916252
38	6377.0	6546.659512
39	29575.5	27158.632753
40	8916.5	6319.811474

```
[192]: # Calculate R-squared for Training data & Display
r_squared_train_for_lasso_cv_hype = round((r2_score(y_train,
    ↪y_train_pred_for_lasso_cv_hype)),4)
print("R-squared for training data: ", r_squared_train_for_lasso_cv_hype)

# Calculate R-squared for Test data & display
r_squared_test_for_lasso_cv_hype = round((r2_score(y_test,
    ↪y_test_pred_for_lasso_cv_hype)),4)
print("R-squared for test data: ", r_squared_test_for_lasso_cv_hype)

# Number of observations (n)
n = len(y_test)

# Number of predictors (k) - you need to replace this with the actual number of
    ↪predictors in your model
k = 19

# Calculate adjusted R Squared for Test data & display
adjusted_r_squared_test_for_lasso_cv_hype = round(1 - ((1 -
    ↪r_squared_test_for_lasso_cv_hype) * (n - 1) / (n - k - 1)),4)

print("Adjusted R-squared for test data:",
    ↪adjusted_r_squared_test_for_lasso_cv_hype)
```

R-squared for training data: 0.9078

R-squared for test data: 0.8617

Adjusted R-squared for test data: 0.7366

```
[193]: # Calculating Mean Squared Error (MSE) for Lasso Regression with
        ↪Cross-Validation Hyperparameter Tuning
mse_lasso_cv_hype = round(mean_squared_error(y_test,
        ↪y_test_pred_for_lasso_cv_hype), 4)

# Calculating Root Mean Squared Error (RSME) for Lasso Regression with
        ↪Cross-Validation Hyperparameter Tuning
rsme_lasso_cv_hype = round(np.sqrt(mean_squared_error(y_test,
        ↪y_test_pred_for_lasso_cv_hype)), 4)

# Calculating Mean Absolute Error (MAE) for Lasso Regression with
        ↪Cross-Validation Hyperparameter Tuning
mae_lasso_cv_hype = round(mean_absolute_error(y_test,
        ↪y_test_pred_for_lasso_cv_hype), 4)

# Calculating Mean Absolute Percentage Error (MAPE) for Lasso Regression with
        ↪Cross-Validation Hyperparameter Tuning
mape_lasso_cv_hype = round(mean_absolute_percentage_error(y_test,
        ↪y_test_pred_for_lasso_cv_hype), 4)

[194]: _# Performance of Lasso Regression Model with Cross Validation & Hyperparameter
        ↪Tuning
print('Performance of Lasso Regression Model')
print("MSE :",round(mean_squared_error(y_test,
        ↪y_test_pred_for_lasso_cv_hype),4))
print("RMSE :",round(np.sqrt(mean_squared_error(y_test,
        ↪y_test_pred_for_lasso_cv_hype)),4))
print('MAE:', round(mean_absolute_error(y_test,
        ↪y_test_pred_for_lasso_cv_hype),4))
print('MAPE:', round(mean_absolute_percentage_error(y_test,
        ↪y_test_pred_for_lasso_cv_hype),4))
```

Performance of Lasso Regression Model

MSE : 6060866.9478

RMSE : 2461.8828


MAE: 2037.7442

MAPE: 0.1865

```
[195]: # Plot the lineplot for Lasso Regression
plt.figure(figsize=(12, 6), dpi = 100)
sns.set_style('darkgrid')
sns.lineplot(data=df_actual_vs_predicted_lasso_model, palette="tab10",
        ↪linewidth=2.5)
plt.title('Actual vs Predicted values for Lasso Regression')
plt.xlabel('Observation')
```



```
plt.ylabel('Values')
plt.show()
```



6.6.3 ML Model - 3 Ridge Regression

Cross- Validation & Hyperparameter Tuning (Ridge Regression)

```
[196]: #Creating an object to run Ridge Regression
ridge = Ridge(max_iter = 6000)

min_alpha = 0.001
max_alpha = 0.2
num_vals = 10

parameters = {'alpha': np.logspace(np.log10(min_alpha), np.log10(max_alpha),
    ↪ num_vals)}

# Applying the GridSearchCV with cross validation and hyperparameter tuning
ridge_cv_hype = GridSearchCV(ridge, parameters, cv = 5, scoring =
    ↪ "neg_mean_squared_error")

# Fit into the Algorithm
ridge_cv_hype.fit(x_train, y_train)
```

```
[196]: GridSearchCV(cv=5, estimator=Ridge(max_iter=6000),
    param_grid={'alpha': array([0.001, 0.00180165, 0.00324594,
    0.00584804, 0.0105361,
    0.01898235, 0.03419952, 0.0616155, 0.11100946, 0.2 ])}),
```

```
scoring='neg_mean_squared_error')
```

```
[197]: # First fit the best alpha value into the model
ridge_alpha = Ridge(alpha = ridge_cv_hype.best_params_['alpha'])
ridge_alpha.fit(x_train, y_train)
```

```
[197]: Ridge(alpha=0.20000000000000004)
```

```
[198]: # Calculate and display the Coefficients
print("Coefficients:", ridge_alpha.coef_)
```

```
Coefficients: [-3.10864887e+03  9.45917393e+02 -3.56752771e+01 -1.48414577e+03
 2.26968954e+01  9.62089678e+02 -1.41723291e+02  4.05842724e+02
 3.37846066e+00 -5.70237242e+03 -8.13420794e+02 -5.49494851e+02
 2.24612310e+01  1.23396744e+02 -2.94129072e+03 -1.41611726e+02
 2.57817639e+01  5.52984298e-03 -1.34482906e+03  8.47244402e+03
 1.80204927e+03]
```

```
[199]: # Calculate and display the Intercept
print("Intercept :", ridge_alpha.intercept_)
```

```
Intercept : 214895.7996472018
```

```
[200]: # Calculate the Predicted values for the Training data
y_train_pred_for_ridge_cv_hype = ridge_alpha.predict(x_train)

# Calculate and Predicted values for Test data
y_test_pred_for_ridge_cv_hype = ridge_alpha.predict(x_test)
```

```
[201]: # Show the Predicted Values for Training Data
print('Predicted Value for Training Data : ',y_train_pred_for_ridge_cv_hype)
```

```
Predicted Value for Training Data : [16011.19261234 27987.46757687
14959.56326337 10705.76517218
 7545.70137128  6922.23401456 14763.77220155 10746.0794175
10237.70670735 22204.33783825  5930.10296356  9124.00947912
10143.9934671  19173.63512634 18872.54121361 10013.17265367
18705.88328554  8792.99210767  6847.90788006  6297.63245063
13787.03569232  7512.75692313 10346.1447515  9430.18052186
14875.93443175  9799.20989541  5991.75839174 10350.14771033
 6317.72420053  6695.80119196  8656.62645722  6770.66814512
 6146.37875269  6473.35438046 17198.45393738  6311.06098045
16849.64652941  7135.51564897 17573.01696886 20301.03012997
20677.17327014  9726.16559008 17012.63860114 10024.86368584
28961.11964256 13889.49877719 14955.83194368 18340.76185671
 6752.84042475 10346.1447515  16972.7333882 10143.9934671
 9187.68586238 17198.15952363  9415.32471429 12151.28449105
17715.9479957  29596.197609  16826.08898221  8646.49107524]
```

```

5321.15827295 6374.82765984 11254.59570913 6273.80421196
6097.70457729 6396.63750574 19175.47772051 6136.12556298
13819.25348991 6243.49176727 17198.45393738 6423.95660548
10396.82166138 8233.40351474 10902.51825448 24017.49186319
32620.88453538 28097.43506738 9664.07146905 18398.37002592
17695.85877922 6625.72390548 9524.42178184 9217.09491263
14384.64971736 6001.76728119 23669.76727175 7205.52237085
9286.46353393 10405.78836786 21112.94882966 20958.28598065
25088.83768237 14780.66450484 6656.99820733 6649.23112911
8769.08558686 6893.79290553 6132.00256553 10409.37792438
12151.28449105 9174.21195991 15709.27478701 7219.39750977
9447.07282516 14806.86773715 13819.25348991 6903.54078176
9242.25560218 9909.48476883 20201.56258985 13449.18962643
10493.45600499 29490.08608302 9570.87477023 13574.19267081
6746.1503419 7429.82742762 29109.77191156 7698.95033806
9145.50125132 17758.8323051 6050.36080613 8556.08147374
14041.39584765 28961.11964256 19473.93975376 6166.58566884
14054.90969029 6841.39240715 19037.52645953 9300.92071693
14489.95904183 9360.78966842 14479.91595421 10344.803736
9163.87529433 17035.46186565 9010.24572388 17278.79111388
26290.09472942 7002.81473132 6922.23401456 10287.5422448
9403.51175458 19918.66611599 9159.42584832 16911.51284462
9668.5377274 11434.07906472 7649.12950688 11018.1526108
26946.65964265 11018.1526108 7589.69407983 5900.17632791
19380.92903744 14845.19725689 17414.89188617 24131.63305875
10838.15106226 18387.12490825 28097.43506738 15588.45798353]

```

```

[202]: # Show the Predicted values for Test data
print('Predicted Value for Test Data :', y_test_pred_for_ridge_cv_hype)

```

```

Predicted Value for Test Data : [ 6787.56044842 19446.81978593 13969.27418002
1937.21997807
10607.49684597 14524.51572188 6263.58390872 7579.80997001
19296.40123121 7644.89496752 17510.04344298 28805.96047373
9430.18052186 13054.00897056 6082.18806395 14092.07275754
13442.89799158 21329.91665313 9391.81635531 5965.89928966
10771.10673896 16990.06807482 12686.29773353 14642.76184494
20599.29701389 5209.79367395 6566.47884465 17705.90516804
6692.02813289 6034.02720642 10014.72830386 12365.1239224
18768.23300226 10004.06973787 6199.57177871 25522.79189062
9851.04230996 14252.88975166 6551.56717337 27985.2417003
5963.94194089]

```

```

[203]: # Create a DataFrame to display the Actual vs Predicted values for the Test
dataset for Ridge Regression
df_actual_vs_predicted_ridge_model = pd.DataFrame({'Actual Values': y_test,
'Predicted Values': y_test_pred_for_ridge_cv_hype})

```

```
df_actual_vs_predicted_ridge_model
```

```
[203]:
```

	Actual Values	Predicted Values
0	6795.0	6787.560448
1	15750.0	19446.819786
2	15250.0	13969.274180
3	5151.0	1937.219978
4	9995.0	10607.496846
5	11199.0	14524.515722
6	5389.0	6263.583909
7	7898.0	7579.809970
8	17199.0	19296.401231
9	6529.0	7644.894968
10	20970.0	17510.043443
11	29575.5	28805.960474
12	10945.0	9430.180522
13	18344.0	13054.008971
14	8916.5	6082.188064
15	9989.0	14092.072758
16	9295.0	13442.897992
17	18920.0	21329.916653
18	7895.0	9391.816355
19	6488.0	5965.899290
20	9959.0	10771.106739
21	15580.0	16990.068075
22	9895.0	12686.297734
23	11549.0	14642.761845
24	15998.0	20599.297014
25	5118.0	5209.793674
26	6938.0	6566.478845
27	16695.0	17705.905168
28	8358.0	6692.028133
29	5499.0	6034.027206
30	7975.0	10014.728304
31	12290.0	12365.123922
32	22018.0	18768.233002
33	8948.0	10004.069738
34	6849.0	6199.571779
35	29575.5	25522.791891
36	11595.0	9851.042310
37	18150.0	14252.889752
38	6377.0	6551.567173
39	29575.5	27985.241700
40	8916.5	5963.941941

```
[204]: # Calculate R-squared for Training data & display
```

```

r_squared_train_for_ridge_cv_hype= round(r2_score(y_train,
↳y_train_pred_for_ridge_cv_hype), 4)
print("R-squared for training data: ", r_squared_train_for_ridge_cv_hype)

# Calculate R-squared for Test data & display
r_squared_test_for_ridge_cv_hype = round(r2_score(y_test,
↳y_test_pred_for_ridge_cv_hype), 4)
print("R-squared for test data: ", r_squared_test_for_ridge_cv_hype)

# Number of observations (n)
n = len(y_test)

# Number of predictors (k) - you need to replace this with the actual number of
↳predictors in your model
k = 19

# Calculate adjusted R Squared for Test data & display
adjusted_r_squared_test_for_ridge_cv_hype = round(1 - ((1 -
↳r_squared_test_for_ridge_cv_hype) * (n - 1) / (n - k - 1)),4)

print("Adjusted R-squared for test data:",
↳adjusted_r_squared_test_for_ridge_cv_hype)

```

R-squared for training data: 0.9092

R-squared for test data: 0.864

Adjusted R-squared for test data: 0.741

[205]:

```

# Calculating Mean Squared Error (MSE) for Ridge Regression with
↳Cross-Validation Hyperparameter Tuning
mse_ridge_cv_hype = round(mean_squared_error(y_test,
↳y_test_pred_for_ridge_cv_hype), 4)

# Calculating Root Mean Squared Error (RSME) for Ridge Regression with
↳Cross-Validation Hyperparameter Tuning
rsme_ridge_cv_hype = round(math.sqrt(mean_squared_error(y_test,
↳y_test_pred_for_ridge_cv_hype))), 4)

# Calculating Mean Absolute Error (MAE) for Ridge Regression with
↳Cross-Validation Hyperparameter Tuning
mae_ridge_cv_hype = round(mean_absolute_error(y_test,
↳y_test_pred_for_ridge_cv_hype), 4)

# Calculating Mean Absolute Percentage Error (MAPE) for Ridge Regression with
↳Cross-Validation Hyperparameter Tuning
mape_ridge_cv_hype = mean_absolute_percentage_error(y_test,
↳y_test_pred_for_ridge_cv_hype)

```

```
[206]: # Performance of Ridge Regression Model
print('Performance of Ridger Regression Model with Cross Validation &
      ↪Hyperparameter Tunning')
print("MSE :",round(mean_squared_error(y_test,
      ↪y_test_pred_for_ridge_cv_hype),4))
print("RMSE :",round(math.sqrt(mean_squared_error(y_test,
      ↪y_test_pred_for_ridge_cv_hype)),4))
print('MAE:', round(mean_absolute_error(y_test,
      ↪y_test_pred_for_ridge_cv_hype),4))
print('MAPE:', round(mean_absolute_percentage_error(y_test,
      ↪y_test_pred_for_ridge_cv_hype),4))
```

Performance of Ridger Regression Model with Cross Validation & Hyperparameter Tunning


MSE : 5961545.047

RMSE : 2441.6275

MAE: 1974.5439

MAPE: 0.1698

```
[207]: # Plot the lineplot for Ridge Regression
plt.figure(figsize=(12, 6), dpi = 100)
sns.set_style('darkgrid')
sns.lineplot(data=df_actual_vs_predicted_ridge_model, palette="tab10",
      ↪linewidth=2.5)
plt.title('Actual vs Predicted values for Ridge Regression')
plt.xlabel('Observations')
plt.ylabel('Values')
plt.show()
```



6.6.4 ML Model - 4 ElasticNet Regression

Cross- Validation & Hyperparameter Tuning (ElasticNet Regression)

```
[208]: # Creating a variable to run Elastic Net Regression
elastic_net = ElasticNet(alpha=0.2, l1_ratio=0.5, max_iter = 6000)

min_alpha = 0.001
max_alpha = 0.2
num_vals = 10

parameters = {'alpha': np.logspace(np.log10(min_alpha), np.log10(max_alpha),
    ↪ num_vals), 'l1_ratio': [0.1, 0.3, 0.5, 0.7, 0.9]}

# Applying GridSearchCV with cross-validation and hyperparameter tuning
elastic_net_cv_hype = GridSearchCV(elastic_net, parameters, cv=5,
    ↪ scoring="neg_mean_squared_error")

# Fit the model
elastic_net_cv_hype.fit(x_train, y_train)
```

```
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 3.084e+08, tolerance: 6.361e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 2.199e+08, tolerance: 6.178e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 2.388e+08, tolerance: 5.247e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 2.545e+08, tolerance: 5.872e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
```

```

Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 2.752e+08, tolerance: 5.900e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 3.065e+08, tolerance: 6.361e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 2.171e+08, tolerance: 6.178e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 2.363e+08, tolerance: 5.247e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 2.533e+08, tolerance: 5.872e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 2.730e+08, tolerance: 5.900e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 3.044e+08, tolerance: 6.361e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 2.140e+08, tolerance: 6.178e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:

```



```

Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 2.336e+08, tolerance: 5.247e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 2.520e+08, tolerance: 5.872e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 2.707e+08, tolerance: 5.900e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 3.021e+08, tolerance: 6.361e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 2.103e+08, tolerance: 6.178e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 2.308e+08, tolerance: 5.247e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 2.507e+08, tolerance: 5.872e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 2.681e+08, tolerance: 5.900e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:

```

```

Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 2.995e+08, tolerance: 6.361e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 2.060e+08, tolerance: 6.178e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 2.277e+08, tolerance: 5.247e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 2.492e+08, tolerance: 5.872e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 2.651e+08, tolerance: 5.900e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 3.145e+08, tolerance: 6.361e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 2.280e+08, tolerance: 6.178e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 2.470e+08, tolerance: 5.247e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:

```

```

Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 2.586e+08, tolerance: 5.872e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 2.817e+08, tolerance: 5.900e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 3.116e+08, tolerance: 6.361e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 2.243e+08, tolerance: 6.178e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 2.431e+08, tolerance: 5.247e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 2.566e+08, tolerance: 5.872e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 2.786e+08, tolerance: 5.900e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 3.084e+08, tolerance: 6.361e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:

```

```

Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 2.199e+08, tolerance: 6.178e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 2.388e+08, tolerance: 5.247e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 2.546e+08, tolerance: 5.872e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 2.752e+08, tolerance: 5.900e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 3.048e+08, tolerance: 6.361e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 2.147e+08, tolerance: 6.178e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 2.342e+08, tolerance: 5.247e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 2.523e+08, tolerance: 5.872e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:

```

```

Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 2.712e+08, tolerance: 5.900e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 3.006e+08, tolerance: 6.361e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 2.078e+08, tolerance: 6.178e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 2.290e+08, tolerance: 5.247e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 2.498e+08, tolerance: 5.872e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 2.664e+08, tolerance: 5.900e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 3.232e+08, tolerance: 6.361e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 2.389e+08, tolerance: 6.178e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:

```

```

Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 2.593e+08, tolerance: 5.247e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 2.647e+08, tolerance: 5.872e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 2.908e+08, tolerance: 5.900e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 3.191e+08, tolerance: 6.361e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 2.339e+08, tolerance: 6.178e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 2.535e+08, tolerance: 5.247e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 2.618e+08, tolerance: 5.872e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 2.866e+08, tolerance: 5.900e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:

```

```

Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 3.145e+08, tolerance: 6.361e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 2.280e+08, tolerance: 6.178e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 2.470e+08, tolerance: 5.247e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 2.586e+08, tolerance: 5.872e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 2.817e+08, tolerance: 5.900e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 3.091e+08, tolerance: 6.361e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 2.209e+08, tolerance: 6.178e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 2.397e+08, tolerance: 5.247e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:

```

```

Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 2.550e+08, tolerance: 5.872e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 2.759e+08, tolerance: 5.900e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 3.024e+08, tolerance: 6.361e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 2.108e+08, tolerance: 6.178e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 2.312e+08, tolerance: 5.247e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 2.508e+08, tolerance: 5.872e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 2.684e+08, tolerance: 5.900e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 3.353e+08, tolerance: 6.361e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:

```



```

Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 2.531e+08, tolerance: 6.178e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 2.767e+08, tolerance: 5.247e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 2.735e+08, tolerance: 5.872e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 3.033e+08, tolerance: 5.900e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 3.297e+08, tolerance: 6.361e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 2.466e+08, tolerance: 6.178e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 2.686e+08, tolerance: 5.247e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 2.694e+08, tolerance: 5.872e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:

```

```

Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 2.975e+08, tolerance: 5.900e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 3.232e+08, tolerance: 6.361e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 2.389e+08, tolerance: 6.178e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 2.593e+08, tolerance: 5.247e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 2.647e+08, tolerance: 5.872e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 2.908e+08, tolerance: 5.900e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 3.155e+08, tolerance: 6.361e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 2.293e+08, tolerance: 6.178e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:

```

```

Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 2.484e+08, tolerance: 5.247e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 2.593e+08, tolerance: 5.872e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 2.827e+08, tolerance: 5.900e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 3.053e+08, tolerance: 6.361e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 2.154e+08, tolerance: 6.178e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 2.348e+08, tolerance: 5.247e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 2.526e+08, tolerance: 5.872e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 2.717e+08, tolerance: 5.900e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:

```

```

Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 3.514e+08, tolerance: 6.361e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 2.714e+08, tolerance: 6.178e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 2.997e+08, tolerance: 5.247e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 2.852e+08, tolerance: 5.872e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 3.196e+08, tolerance: 5.900e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 3.440e+08, tolerance: 6.361e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 2.631e+08, tolerance: 6.178e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 2.892e+08, tolerance: 5.247e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:

```

```

Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 2.798e+08, tolerance: 5.872e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 3.122e+08, tolerance: 5.900e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 3.353e+08, tolerance: 6.361e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 2.531e+08, tolerance: 6.178e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 2.767e+08, tolerance: 5.247e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 2.735e+08, tolerance: 5.872e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 3.033e+08, tolerance: 5.900e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 3.246e+08, tolerance: 6.361e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:

```

```

Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 2.406e+08, tolerance: 6.178e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 2.613e+08, tolerance: 5.247e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 2.657e+08, tolerance: 5.872e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 2.923e+08, tolerance: 5.900e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 3.098e+08, tolerance: 6.361e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 2.219e+08, tolerance: 6.178e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 2.407e+08, tolerance: 5.247e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 2.555e+08, tolerance: 5.872e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:

```

```

Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 2.767e+08, tolerance: 5.900e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 3.713e+08, tolerance: 6.361e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 2.939e+08, tolerance: 6.178e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 3.276e+08, tolerance: 5.247e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 2.996e+08, tolerance: 5.872e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 3.397e+08, tolerance: 5.900e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 3.623e+08, tolerance: 6.361e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 2.838e+08, tolerance: 6.178e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:

```

```

Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 3.151e+08, tolerance: 5.247e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 2.931e+08, tolerance: 5.872e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 3.307e+08, tolerance: 5.900e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 3.514e+08, tolerance: 6.361e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 2.714e+08, tolerance: 6.178e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 2.997e+08, tolerance: 5.247e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 2.852e+08, tolerance: 5.872e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 3.196e+08, tolerance: 5.900e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:

```



```

Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 3.372e+08, tolerance: 6.361e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 2.553e+08, tolerance: 6.178e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 2.795e+08, tolerance: 5.247e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 2.749e+08, tolerance: 5.872e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 3.053e+08, tolerance: 5.900e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 3.165e+08, tolerance: 6.361e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 2.307e+08, tolerance: 6.178e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 2.499e+08, tolerance: 5.247e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:

```

```

Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 2.600e+08, tolerance: 5.872e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 2.839e+08, tolerance: 5.900e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 3.914e+08, tolerance: 6.361e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 3.185e+08, tolerance: 6.178e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 3.578e+08, tolerance: 5.247e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 3.149e+08, tolerance: 5.872e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 3.618e+08, tolerance: 5.900e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 3.823e+08, tolerance: 6.361e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:

```

```

Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 3.076e+08, tolerance: 6.178e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 3.446e+08, tolerance: 5.247e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 3.081e+08, tolerance: 5.872e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 3.519e+08, tolerance: 5.900e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 3.706e+08, tolerance: 6.361e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 2.936e+08, tolerance: 6.178e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 3.274e+08, tolerance: 5.247e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 2.993e+08, tolerance: 5.872e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:

```

```

Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 3.394e+08, tolerance: 5.900e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 3.536e+08, tolerance: 6.361e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 2.741e+08, tolerance: 6.178e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 3.031e+08, tolerance: 5.247e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 2.869e+08, tolerance: 5.872e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 3.219e+08, tolerance: 5.900e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 3.261e+08, tolerance: 6.361e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 2.424e+08, tolerance: 6.178e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:

```

```

Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 2.635e+08, tolerance: 5.247e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 2.668e+08, tolerance: 5.872e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 2.939e+08, tolerance: 5.900e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 1.940e+07, tolerance: 6.361e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 1.192e+08, tolerance: 6.178e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 1.239e+08, tolerance: 5.247e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 1.373e+08, tolerance: 5.872e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 1.297e+08, tolerance: 5.900e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:

```

```

Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 2.031e+08, tolerance: 6.361e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 2.567e+08, tolerance: 6.178e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 2.891e+08, tolerance: 5.247e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 2.523e+08, tolerance: 5.872e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 2.907e+08, tolerance: 5.900e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 3.681e+08, tolerance: 6.361e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 3.092e+08, tolerance: 6.178e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 3.484e+08, tolerance: 5.247e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:

```

```

Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 2.883e+08, tolerance: 5.872e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 3.512e+08, tolerance: 5.900e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 3.715e+08, tolerance: 6.361e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 2.960e+08, tolerance: 6.178e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 3.305e+08, tolerance: 5.247e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 3.006e+08, tolerance: 5.872e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 3.412e+08, tolerance: 5.900e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 3.391e+08, tolerance: 6.361e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:

```

```

Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 2.576e+08, tolerance: 6.178e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 2.824e+08, tolerance: 5.247e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 2.763e+08, tolerance: 5.872e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 3.073e+08, tolerance: 5.900e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 6.161e+06, tolerance: 6.178e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 5.926e+06, tolerance: 5.247e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 8.984e+06, tolerance: 5.872e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 6.574e+06, tolerance: 5.900e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:

```



```

Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 2.959e+08, tolerance: 6.361e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 2.796e+08, tolerance: 6.178e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 3.184e+08, tolerance: 5.247e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 1.371e+08, tolerance: 5.872e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 3.174e+08, tolerance: 5.900e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 3.553e+08, tolerance: 6.361e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 2.763e+08, tolerance: 6.178e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 3.060e+08, tolerance: 5.247e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:

```

```

Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 2.881e+08, tolerance: 5.872e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 3.237e+08, tolerance: 5.900e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 3.599e+08, tolerance: 6.361e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 2.928e+08, tolerance: 6.178e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 3.283e+08, tolerance: 5.247e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 2.896e+08, tolerance: 5.872e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 3.365e+08, tolerance: 5.900e+05
    model = cd_fast.enet_coordinate_descent(
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 3.913e+08, tolerance: 7.399e+05
    model = cd_fast.enet_coordinate_descent(

```

```
[208]: GridSearchCV(cv=5, estimator=ElasticNet(alpha=0.2, max_iter=6000),
                param_grid={'alpha': array([0.001, 0.00180165, 0.00324594,
                0.00584804, 0.0105361,
                0.01898235, 0.03419952, 0.0616155, 0.11100946, 0.2,
                'l1_ratio': [0.1, 0.3, 0.5, 0.7, 0.9]}],
                scoring='neg_mean_squared_error')
```

```
[209]: # Fit the best alpha and l1_ratio values into the model
elastic_net_best = ElasticNet(alpha=elastic_net_cv_hype.best_params_['alpha'],
    ↪ l1_ratio=elastic_net_cv_hype.best_params_['l1_ratio'])
elastic_net_best.fit(x_train, y_train)
```

```
/usr/local/lib/python3.10/dist-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 3.991e+08, tolerance: 7.399e+05
    model = cd_fast.enet_coordinate_descent(
```

```
[209]: ElasticNet(alpha=0.010536102768906645, l1_ratio=0.3)
```

```
[210]: # Calculate and display the Coefficients
print("Coefficients:", elastic_net_best.coef_)
```

```
Coefficients: [-2.73370161e+02  7.50572214e+02 -1.47863047e+02 -4.89806856e+02
 3.05905619e+01  1.22795725e+03 -9.19022925e+01  1.80115114e+02
 3.34158117e+00 -2.11503807e+03 -7.57957630e+02 -5.25132829e+02
 6.61450981e+00  1.98827287e+02 -3.29880694e+03 -1.30518996e+02
 2.49310431e+01  1.11283044e-02 -1.15188476e+03  5.64647341e+03
 1.95155327e+03]
```

```
[211]: # Calculate and display Intercept
print("Intercept:", elastic_net_best.intercept_)
```

```
Intercept: 29594.097929877942
```

```
[212]: # Calculate the Predicted values for the Training data
y_train_pred_for_elastic_net_cv_hype = elastic_net_best.predict(x_train)

# Calculate the Predicted values for the Test data
y_test_pred_for_elastic_net_cv_hype = elastic_net_best.predict(x_test)
```

```
[213]: # Show the predicted values for the training data
print('Predicted Values for Training Data:',
    ↪ y_train_pred_for_elastic_net_cv_hype)
```

```
Predicted Values for Training Data: [16108.69402839 26961.53887765 14550.4523736
11000.828162
```

```

7053.87020118 6453.79029914 15130.04704521 10871.71768714
10327.58529838 21329.74152816 5779.60315853 9297.21419708
10359.34672177 19776.91786155 19546.91054751 10293.93940851
18875.47684479 8961.14021841 6380.2755135 5784.33924265
13568.98824968 7286.77902784 10679.72616882 9315.59330997
15506.74047688 9867.40972249 5483.11267285 10554.59582059
6214.23856533 6701.27506591 8879.37459559 6795.87318143
6043.64743049 6851.88681975 17495.01456409 6257.15762151
16808.16541209 7421.21258782 17065.31436993 21412.90895386
20928.3836879 9151.3363035 17311.22759998 10117.06568495
27249.71860785 13766.79558291 15153.06774439 18071.84647017
6977.85567081 10679.72616882 17427.05250429 10359.34672177
9551.90253744 17557.278643 8754.08250775 12138.0014842
17060.99199574 28245.02284694 16703.24741942 8869.34985209
5353.74329674 6248.53271065 11205.84871326 6170.79801018
5594.83418595 6857.64481535 19433.27937989 5965.30301528
13418.01721009 6190.3259982 17495.01456409 6297.80729521
10729.8498863 8432.02800747 10759.48446338 24009.40635313
33310.83538556 27469.25644213 9733.74647587 18523.54851086
18278.66693832 6999.05567888 8963.1528439 8709.23319387
14727.22095896 5506.33855109 24397.56952345 7048.45325992
9696.51345104 10760.29644206 21293.57337201 19965.58054445
25059.99192462 15146.75495104 6179.42246087 6932.07261653
8993.88227802 6952.33544208 6080.05381974 10871.512887
12138.0014842 8803.672248 16288.42632737 7067.27683812
9332.3012158 15343.21843941 13418.01721009 6400.91003122
9414.16953787 10181.37614092 21449.04707086 13234.83013313
10888.59279704 29023.21111257 9744.08988137 13358.46863625
6742.16139825 7742.52150065 27396.74817913 7419.98459367
9542.78686735 17249.10133404 6120.32867676 8161.80918956
14470.31136694 27249.71860785 20555.75114939 5881.27946025
14483.6776916 7166.01018434 18742.85467844 9447.08646034
14792.22330482 9770.02823668 14847.50675028 10741.56322575
8889.93887213 16991.95237619 9348.87218946 17244.34204869
27093.86177315 6847.95838999 6453.79029914 10643.34110126
9321.31234202 19997.37435811 9675.44603291 16807.01117357
9739.0663307 11366.70781931 7669.05295048 11232.43796796
25393.7028472 11232.43796796 7615.22189501 5433.85325928
19736.19818253 15421.00918276 17398.22952033 23946.16168281
10848.07796392 18688.86589456 27469.25644213 16065.89445268]

```

```

[214]: # Show the predicted values for the test data
print('Predicted Values for Test Data:', y_test_pred_for_elastic_net_cv_hype)

```

```

Predicted Values for Test Data: [ 6812.58108726 20553.89246256 14059.75716384
-1008.84562722
10771.23924077 14948.15747361 6096.4027577 7814.49926488
20375.17939863 7366.51929502 18094.87997422 31079.87388922

```

```

9315.59330997 13534.28286751 5589.75075478 14520.43508442
13138.14730064 20333.15447266 9566.9860796 6036.78914762
10833.37015878 16930.77125244 12267.83848461 15065.1128144
21727.19232144 5067.91841611 6862.34454864 16880.80201625
6917.70720983 5983.14796594 10107.04094145 12127.36096144
19574.66775227 10495.57763836 6146.88544305 25706.60324497
9662.96626559 14596.8992935 6098.52778126 27490.0138902
5472.79541398]

```

```

[215]: # Create a DataFrame to display the Actual vs Predicted values for the Test
        ↳ dataset for Elastic Net Regression
df_actual_vs_predicted_elastic_net_model = pd.DataFrame({'Actual Values':
        ↳ y_test, 'Predicted Values': y_test_pred_for_elastic_net_cv_hype})
df_actual_vs_predicted_elastic_net_model

```

```

[215]:
Actual Values Predicted Values
0          6795.0      6812.581087
1         15750.0     20553.892463
2         15250.0     14059.757164
3          5151.0     -1008.845627
4          9995.0     10771.239241
5         11199.0     14948.157474
6          5389.0      6096.402758
7          7898.0      7814.499265
8         17199.0     20375.179399
9          6529.0      7366.519295
10        20970.0     18094.879974
11        29575.5     31079.873889
12        10945.0      9315.593310
13        18344.0     13534.282868
14          8916.5      5589.750755
15          9989.0     14520.435084
16          9295.0     13138.147301
17        18920.0     20333.154473
18          7895.0      9566.986080
19          6488.0      6036.789148
20          9959.0     10833.370159
21        15580.0     16930.771252
22          9895.0     12267.838485
23        11549.0     15065.112814
24        15998.0     21727.192321
25          5118.0      5067.918416
26          6938.0      6862.344549
27        16695.0     16880.802016
28          8358.0      6917.707210
29          5499.0      5983.147966
30          7975.0     10107.040941

```

31	12290.0	12127.360961
32	22018.0	19574.667752
33	8948.0	10495.577638
34	6849.0	6146.885443
35	29575.5	25706.603245
36	11595.0	9662.966266
37	18150.0	14596.899294
38	6377.0	6098.527781
39	29575.5	27490.013890
40	8916.5	5472.795414

```
[216]: # Calculate R-squared for Training data
r_squared_train_for_elastic_net_cv_hype = round(r2_score(y_train,
    ↳y_train_pred_for_elastic_net_cv_hype), 4)
print("R-squared for training data:", r_squared_train_for_elastic_net_cv_hype)

# Calculate R-squared for Test data
r_squared_test_for_elastic_net_cv_hype = round(r2_score(y_test,
    ↳y_test_pred_for_elastic_net_cv_hype), 4)
print("R-squared for test data:", r_squared_test_for_elastic_net_cv_hype)

# Number of observations (n)
n = len(y_test)

# Number of predictors (k) - you need to replace this with the actual number of
    ↳predictors in your model
k = 19

# Calculate adjusted R Squared for Test data & display
adjusted_r_squared_test_for_elastic_net_cv_hype = round(1 - ((1 -
    ↳r_squared_test_for_elastic_net_cv_hype) * (n - 1) / (n - k - 1)),4)

print("Adjusted R-squared for test data:",
    ↳adjusted_r_squared_test_for_elastic_net_cv_hype)
```

R-squared for training data: 0.9012

R-squared for test data: 0.837

Adjusted R-squared for test data: 0.6895

```
[217]: # Calculating Mean Squared Error (MSE) for Elastic Net with Cross-Validation
    ↳Hyperparameter Tuning
mse_elastic_net_cv_hype = round(mean_squared_error(y_test,
    ↳y_test_pred_for_elastic_net_cv_hype), 4)

# Calculating Root Mean Squared Error (RSME) for Elastic Net with
    ↳Cross-Validation Hyperparameter Tuning
```

```

rsme_elastic_net_cv_hype = round(math.sqrt(mean_squared_error(y_test,
↳y_test_pred_for_elastic_net_cv_hype)), 4)

# Calculating Mean Absolute Error (MAE) for Elastic Net with Cross-Validation
↳Hyperparameter Tuning
mae_elastic_net_cv_hype = mean_absolute_error(y_test,
↳y_test_pred_for_elastic_net_cv_hype)

# Calculating Mean Absolute Percentage Error (MAPE) for Elastic Net with
↳Cross-Validation Hyperparameter Tuning
mape_elastic_net_cv_hype = round(mean_absolute_percentage_error(y_test,
↳y_test_pred_for_elastic_net_cv_hype), 4)

```

```

[218]: # Performance of Elastic Net Regression Model
print('Performance of Elastic Net Regression Model with Cross Validation &
↳Hyperparameter Tuning')
print("MSE:", round(mean_squared_error(y_test,
↳y_test_pred_for_elastic_net_cv_hype), 4))
print("RMSE:", round(math.sqrt(mean_squared_error(y_test,
↳y_test_pred_for_elastic_net_cv_hype)), 4))
print('MAE:', mean_absolute_error(y_test, y_test_pred_for_elastic_net_cv_hype))
print('MAPE:', round(mean_absolute_percentage_error(y_test,
↳y_test_pred_for_elastic_net_cv_hype), 4))


```

Performance of Elastic Net Regression Model with Cross Validation &
 Hyperparameter Tuning
 MSE: 7143361.1727
 RMSE: 2672.7067
 MAE: 2092.354269739015
 MAPE: 0.1872

```

[219]: # Plot the lineplot to visualize the match between original and predicted values
plt.figure(figsize=(12, 6), dpi = 100)
sns.set_style('darkgrid')
sns.lineplot(data=df_actual_vs_predicted_elastic_net_model, palette="tab10",
↳linewidth=2.5)
plt.title('Actual vs Predicted values for Elastic Net Regression')
plt.xlabel('Observations')
plt.ylabel('Values')
plt.show()

```



6.7 7. Interpretation of Results

```
[220]: # Create a dictionary with model names as keys and metric values as lists
data = {
    'Linear Regression': [r_squared_test_for_linear_regression,
    ↪adjusted_r_squared_test_for_linear_regression,
    ↪r_squared_train_for_linear_regression, mse_lr, rsme_lr, mae_lr, mape_lr],
    'Lasso Regression': [r_squared_test_for_lasso_cv_hype,
    ↪adjusted_r_squared_test_for_lasso_cv_hype, r_squared_train_for_lasso_cv_hype,
    ↪mse_lasso_cv_hype, rsme_lasso_cv_hype, mae_lasso_cv_hype,
    ↪mape_lasso_cv_hype],
    'Ridge Regression': [r_squared_test_for_ridge_cv_hype,
    ↪adjusted_r_squared_test_for_ridge_cv_hype,
    ↪r_squared_train_for_ridge_cv_hype, mse_ridge_cv_hype, rsme_ridge_cv_hype,
    ↪mae_ridge_cv_hype, mape_ridge_cv_hype],
    'Elastic Net Regression': [r_squared_test_for_elastic_net_cv_hype,
    ↪adjusted_r_squared_test_for_elastic_net_cv_hype,
    ↪r_squared_train_for_elastic_net_cv_hype, mse_elastic_net_cv_hype,
    ↪rsme_elastic_net_cv_hype, mae_elastic_net_cv_hype, mape_elastic_net_cv_hype]
}

# Create the DataFrame
model_comparison_for_test_train_data = pd.DataFrame(data, index=['R-squared_
    ↪Test', 'Adjusted R Squared Test', 'R-squared Train ', 'MSE', 'RSME', 'MAE',
    ↪'MAPE'])

# Transpose the DataFrame for a better view
```



```
model_comparison_for_test_train_data = model_comparison_for_test_train_data.T
model_comparison_for_test_train_data
```

[220]:	R-squared Test	Adjusted R Squared Test	\		
Linear Regression	0.8723	0.7568			
Lasso Regression	0.8617	0.7366			
Ridge Regression	0.8640	0.7410			
Elastic Net Regression	0.8370	0.6895			

	R-squared Train	MSE	RSME	MAE	\
Linear Regression	0.9103	5.596513e+06	2365.6951	1967.19520	
Lasso Regression	0.9078	6.060867e+06	2461.8828	2037.74420	
Ridge Regression	0.9092	5.961545e+06	2441.6275	1974.54390	
Elastic Net Regression	0.9012	7.143361e+06	2672.7067	2092.35427	

	MAPE
Linear Regression	0.169800
Lasso Regression	0.186500
Ridge Regression	0.169818
Elastic Net Regression	0.187200

Now we will look at different models simultaneously and compare them different metrics. For this project we will be comparing the results of different models on 5 different metrics i.e

-
1. **R Squared (Test Data)** - R-squared measures the proportion of the variance in the dependent variable that is predictable from the independent variables. It provides an indication of how well the model fits the test data. R-squared ranges from 0 to 1, where 1 indicates a perfect fit. However, a high R-squared doesn't guarantee a good model, so it's crucial to assess performance on unseen data.
-
2. **R Squared (Training Data)** - Similar to R-squared for test data but calculated on the training data. It helps understand how well the model fits the training data. While a high R-squared on training data may indicate a good fit, evaluating the model on test data is essential for generalization.
-
3. **Adjusted R Squared (Test Data)** - Adjusted R-squared is a modified version that penalizes the inclusion of irrelevant predictors. Useful for models with multiple predictors, it considers the number of predictors and penalizes models that add less value, helping to prevent overfitting.
-
4. **Mean Squared Error** - MSE is the average of the squared differences between predicted and actual values. It provides a measure of the average squared deviation between predicted

and actual values. Lower MSE values indicate better model performance.

5. **Root Mean Squared Error** - RMSE is the square root of the mean squared error. Like MSE, it measures the average magnitude of errors. Its units are the same as the dependent variable, making it easier to interpret.
-

6. **Mean Absolute Error** - MAE is the average of the absolute differences between predicted and actual values. It provides a measure of the average absolute deviation between predicted and actual values, giving equal weight to all errors.
-

7. **Mean Absolute Percentage Error** - MAPE is the average percentage difference between predicted and actual values, expressed as a percentage of the actual values. It is useful when the scale of the dependent variable is significant, providing a percentage measure of the average prediction error.

Detailed Analysis Here is a detailed analysis of the results:

Linear Regression

Linear regression still has the best metrics overall. The higher test R-squared of 0.8723 and adjusted R-squared of 0.7568 indicate it generalizes very well while achieving high accuracy with the 0.9103 train R-squared. The MSE, RMSE, MAE and MAPE are the lowest showing it makes the closest price predictions. There is some overfitting evident from the train and test gap, but the model still generalizes reasonably well. Overall linear regression still produces the most accurate and generalizable model.

Lasso Regression

Lasso regression now has slightly improved metrics getting closer to linear regression's performance. The test R-squared of 0.8665 and adjusted R-squared of 0.7457 indicate good generalizability and accuracy. All the error metrics have also improved compared to before. There is still a gap between train and test showing some overfitting. With hyperparameter tuning lasso can potentially match or exceed linear regression.

Ridge Regression

Ridge regression maintains its middle ground between linear and lasso regression. The metrics are better than lasso showing higher accuracy but train and test gap indicates more overfitting than lasso. With tuning, ridge can achieve further optimization between accuracy and preventing overfit.

Elastic Net Regression

Elastic net still has the lowest metrics but has improved over previous results. There is still significant gap between train and test R-squared indicating it focuses heavily on reducing overfit which affects accuracy. Further tuning can help balance accuracy and overfitting better. Overall it achieves reasonable accuracy with the highest generalization capability.

6.8 8. Conclusion

Based on the regression analysis, the linear model provides the most accurate and generalizable relationship between car features and pricing. The high R-squared and adjusted R-squared along with lowest error metrics prove it models the prices effectively based on independent variables like horsepower, dimensions, engine size etc.

Therefore, Geely should optimize the car design and component parameters like horsepower, curb-weight, engine-size etc according to their target base/premium pricing position in the market. The linear model coefficients will tell them exactly how each parameter impacts pricing. Analyzing competitors can reveal the typical value ranges.

Additionally, new data should continually be fed to the linear model to account for changing market dynamics over time and improve accuracy. The regular model re-training will allow Geely to keep pricing dynamic with changing trends.

6.8.1 Hurrah! You have successfully completed your Machine Learning Capstone Project !!!